saama

SEP 2024
MONTHLY MEETUP

AI Agents

presented by
Navaneeth Malingan,
Founder/CEO - Nunnari Labs

About Me

I Build Products and Communities

AI and IoT Researcher, Developer, Educator

Ex-Director - Innovation (KGiSL Edu)

NUNNARI LABS

AI Tamil Nadu

TensorFlow User Group Coimbatore

Navaneeth Malingan,
Founder/CEO - Nunnari Labs

NUNNARI
LABS

Let's start with LLMs!

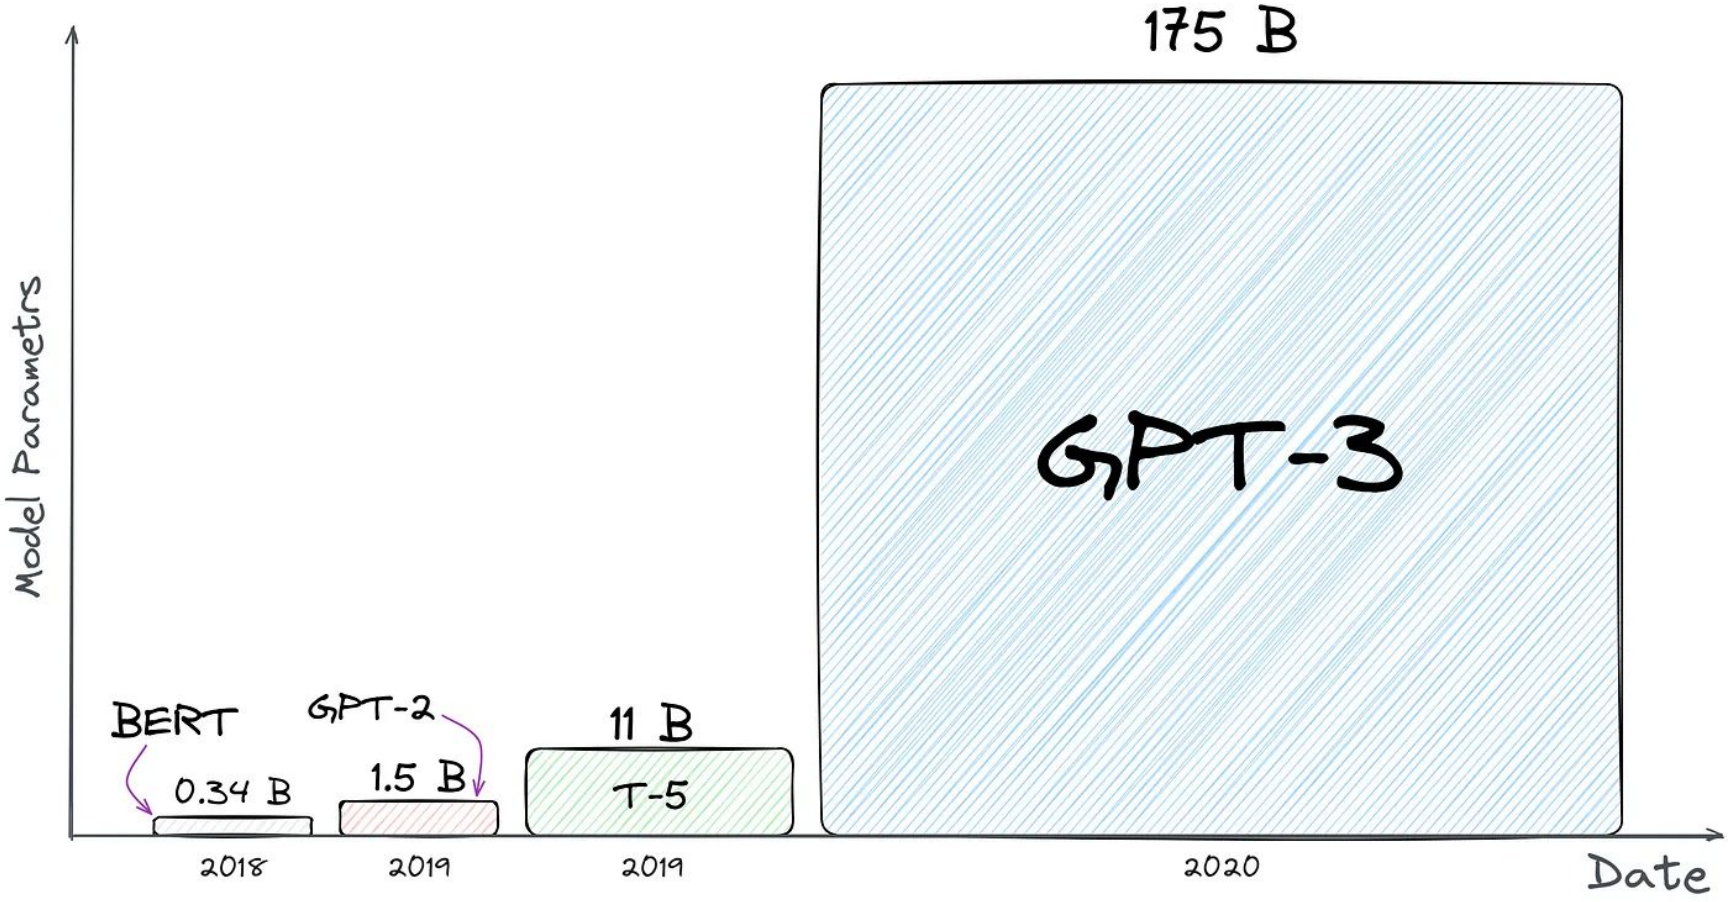

NUNNARI
LABS

LLM Capabilities

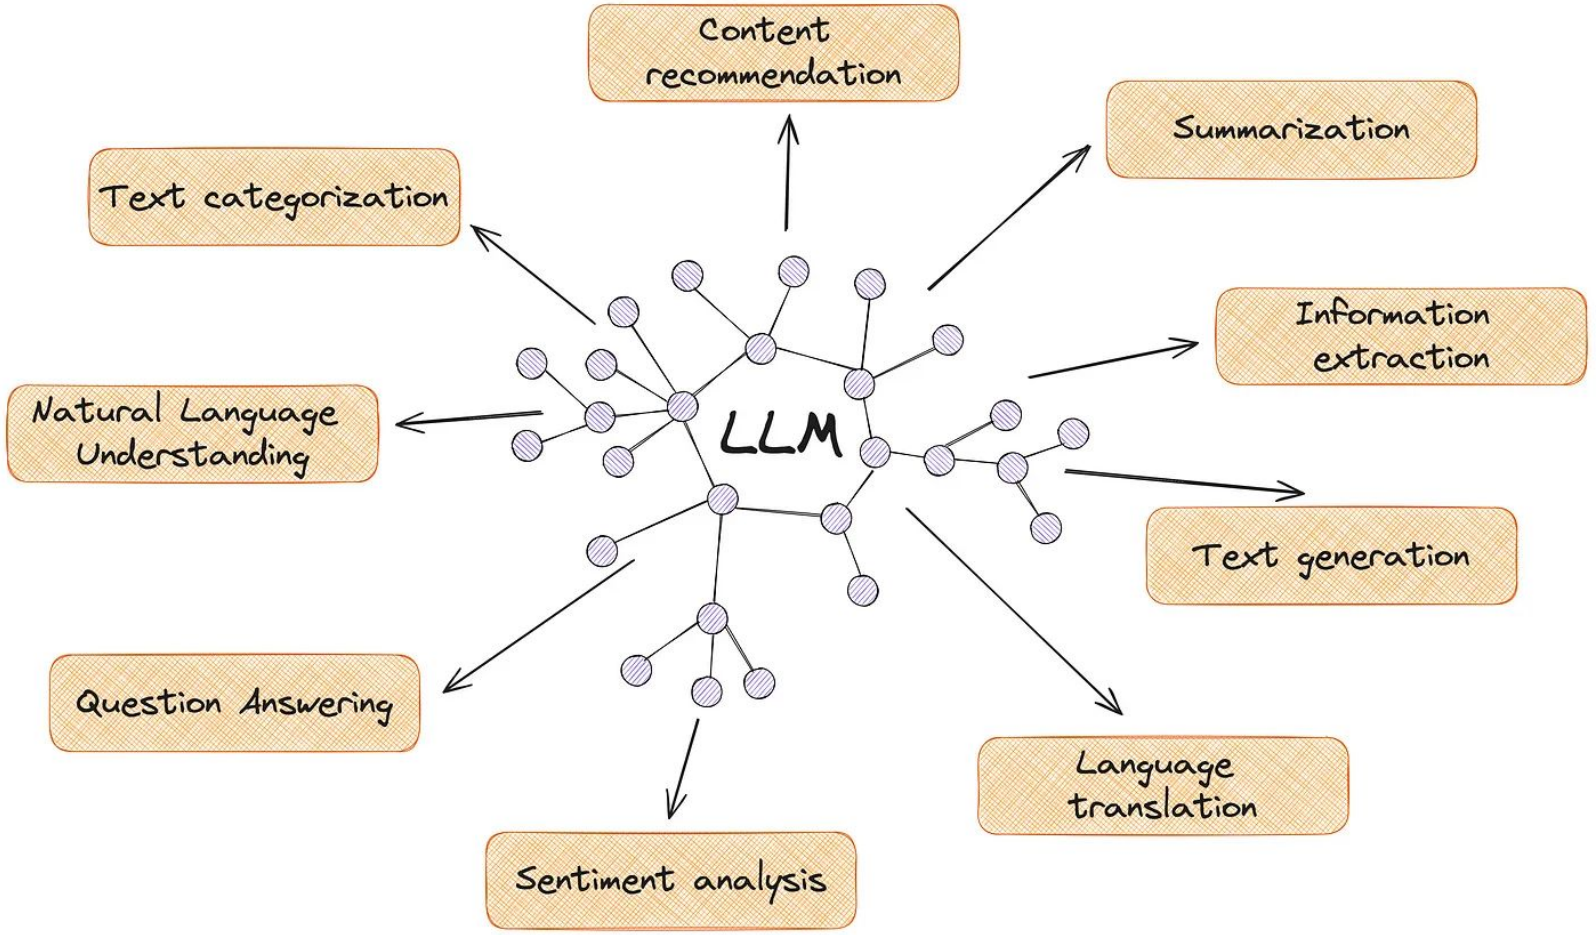

Prompt Engineering

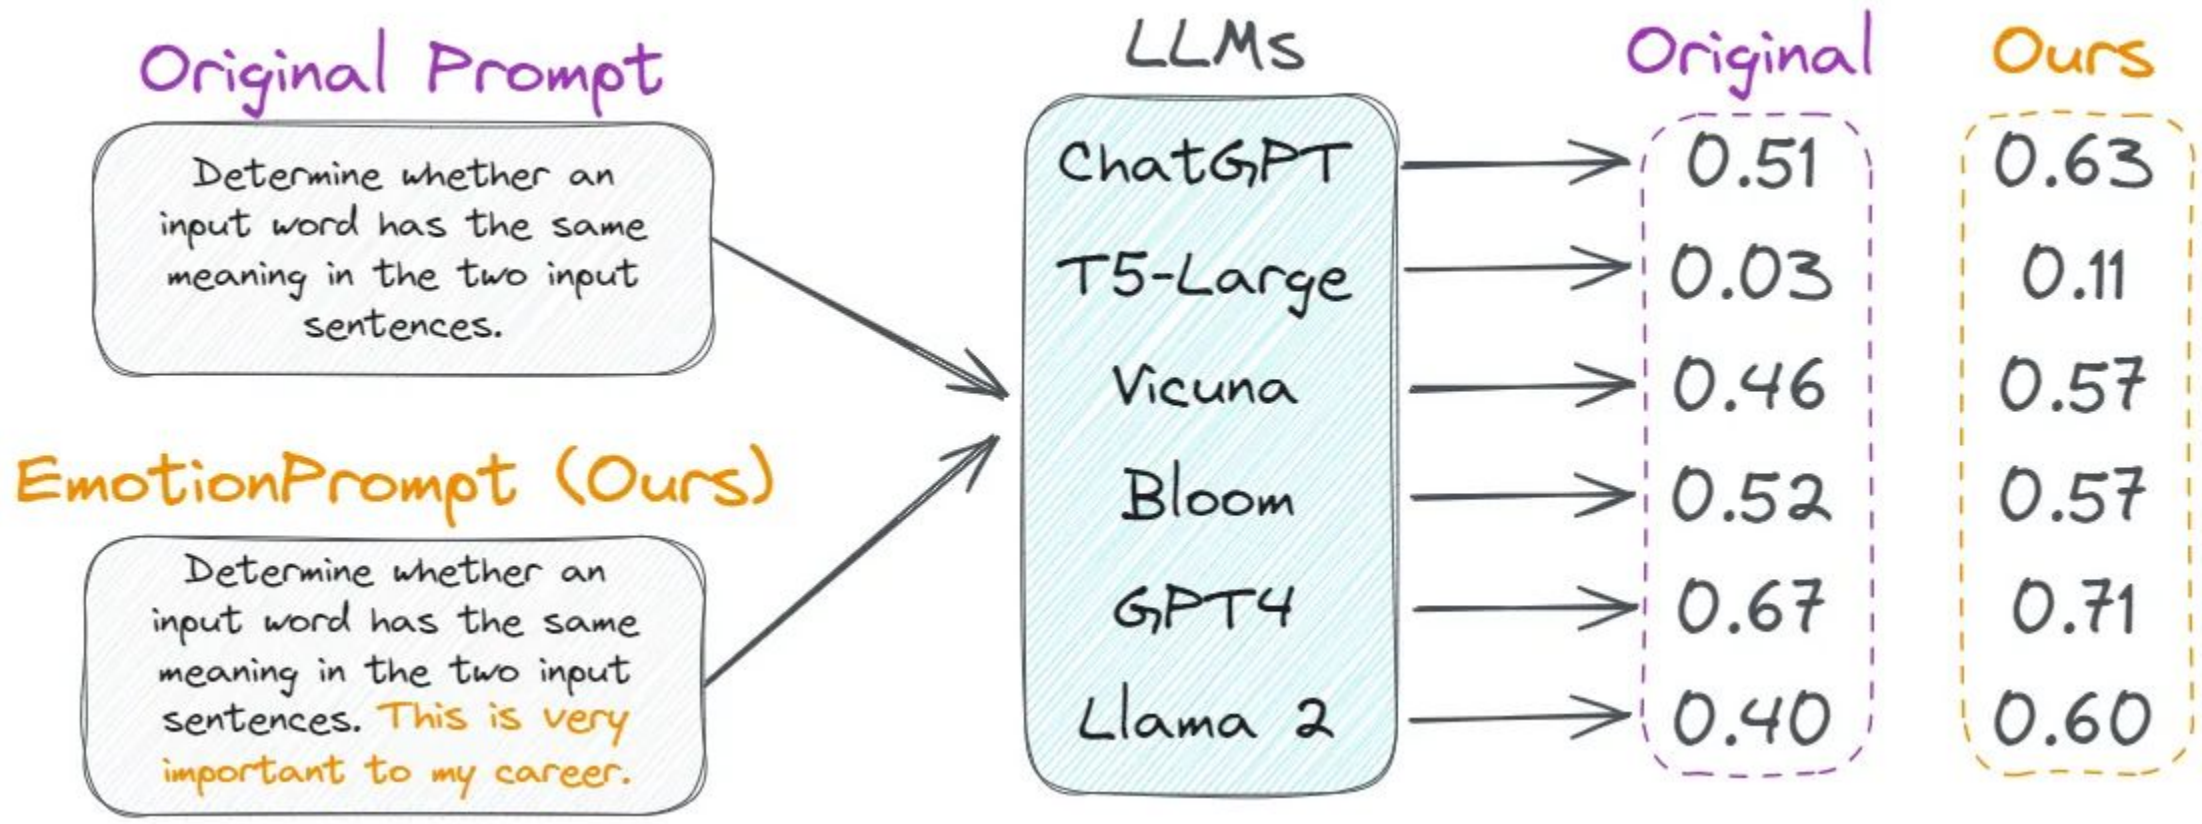

NUNNARI
LABS

Types of Prompting

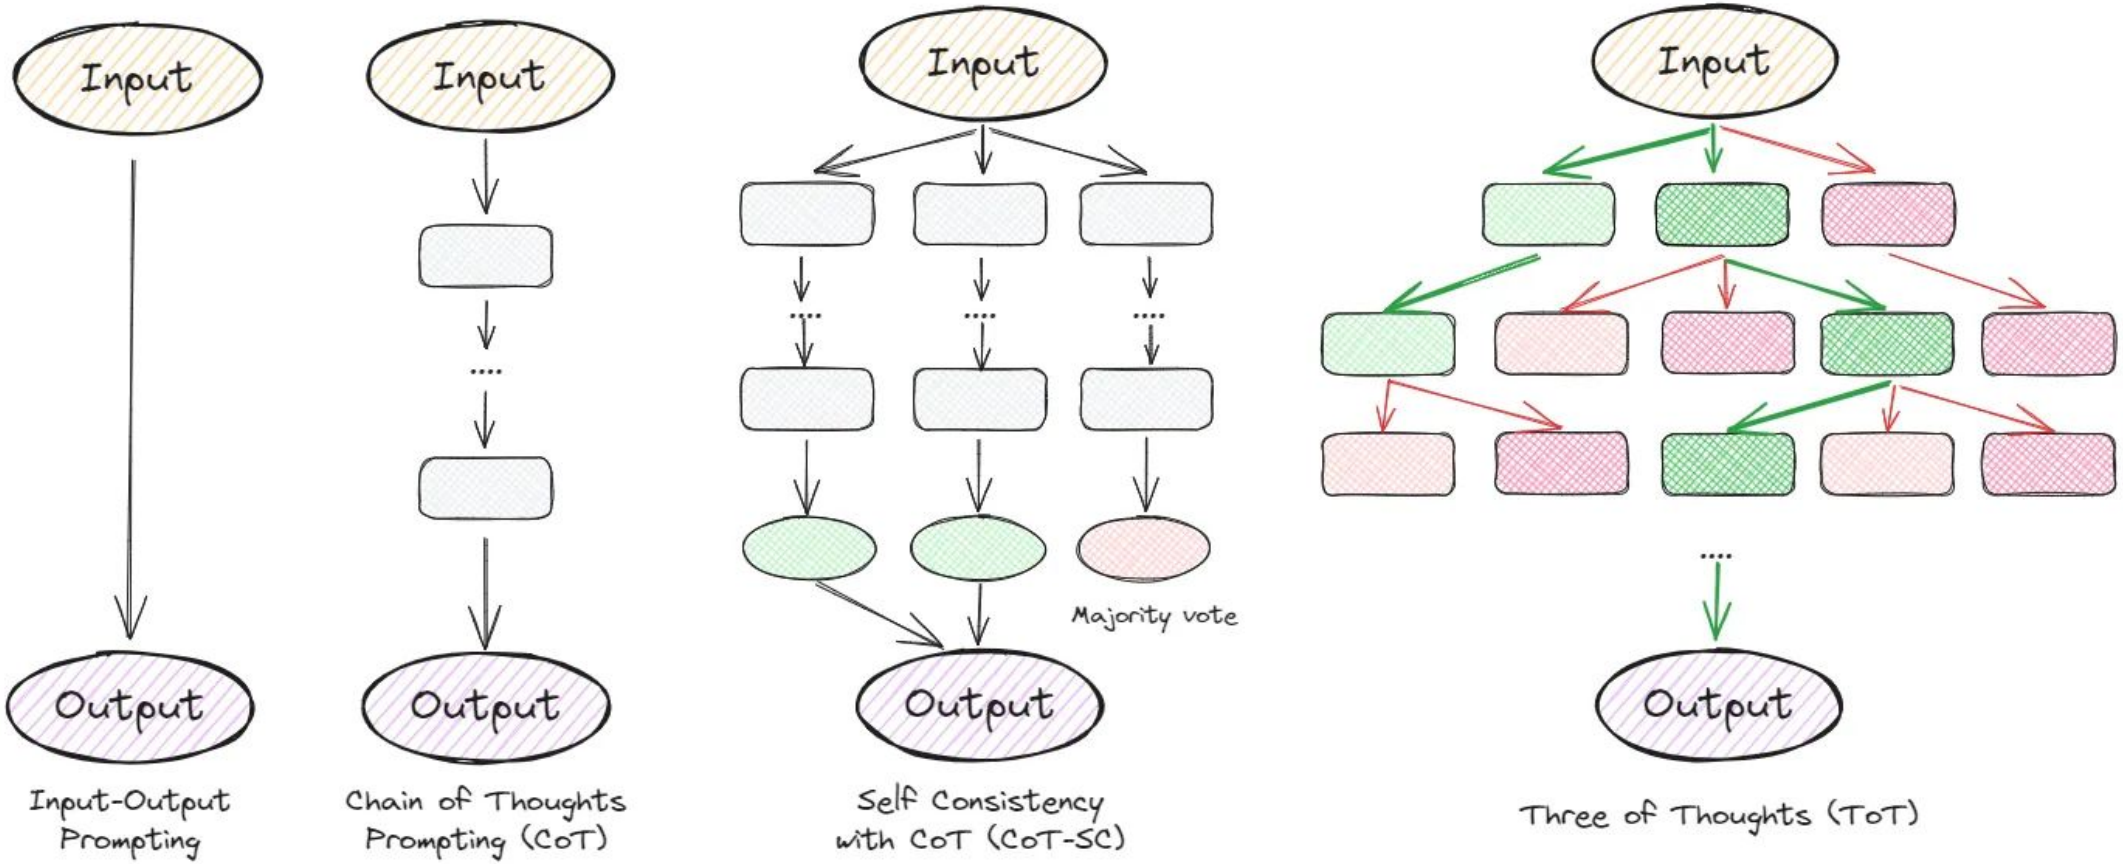

NUNNARI
LABS

Retrieval Augmented Generation

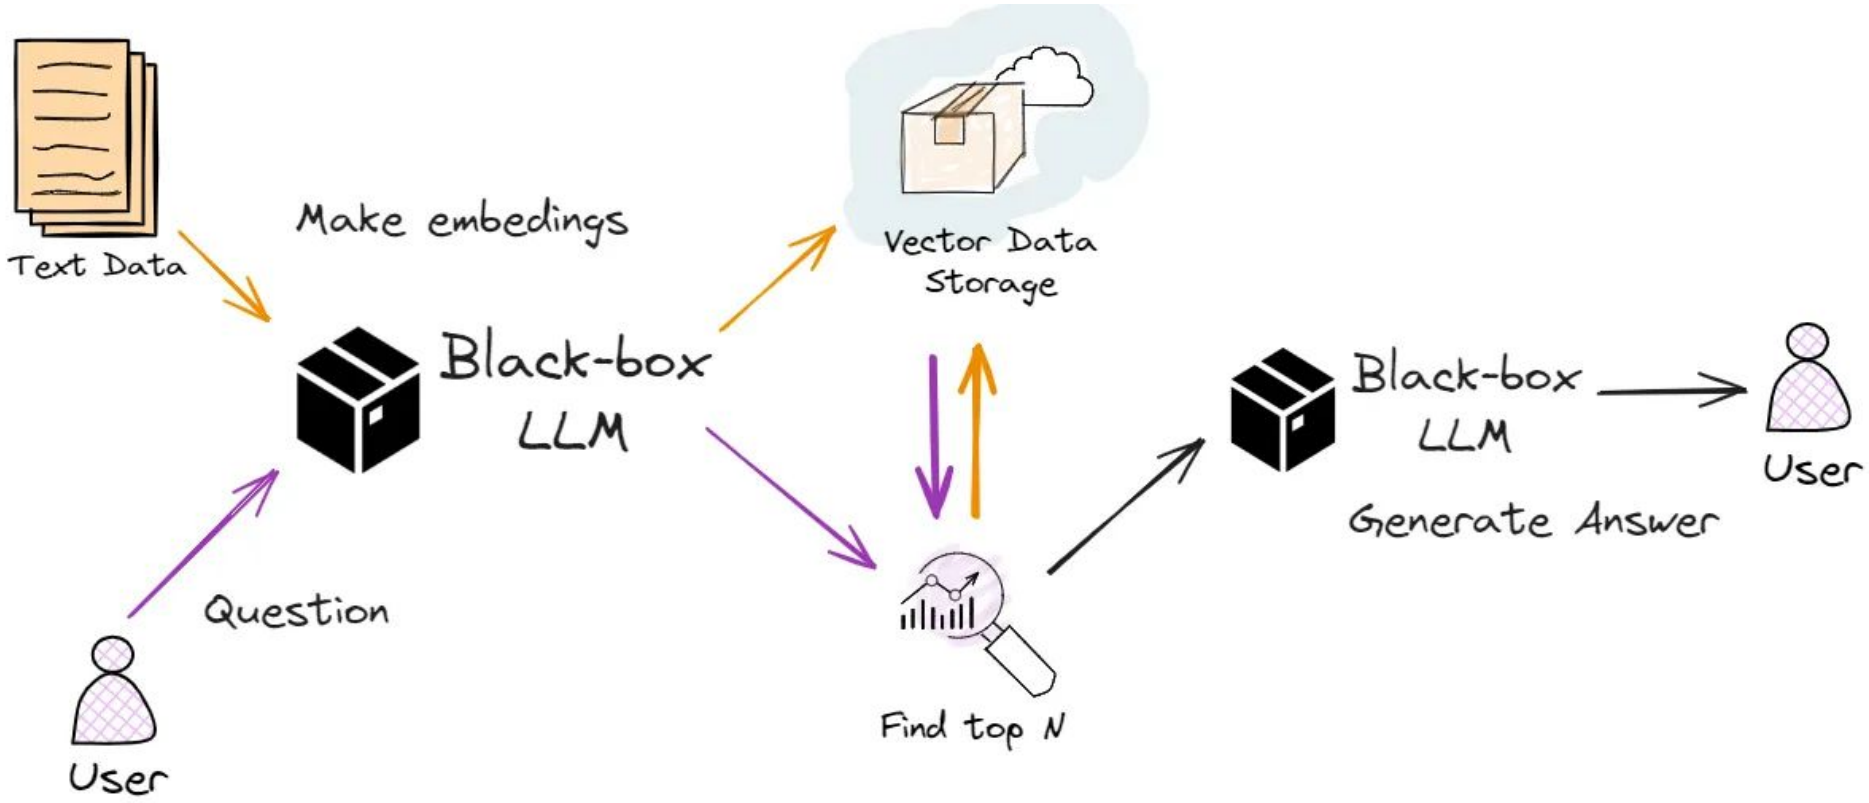

LLM Hallucination

NUNNARI LABS

LLMs hallucinate when their training dataset contains limited, outdated, or contradictory information about the question posed to them.

NUNNARI
LABS

What is a Generative AI Agent?

Navaneeth Malingan

NUNNARI
LABS

A Generative AI Agent is defined as an application that tries to achieve a goal by observing the world and acting upon it using the tools it has at its disposal.

Observe, Act, Achieve!

GenAI Agent Categorization

NUNNARI LABS

Modality Based

- **LLM (Large Language Model)**, focus on language inputs
- **VLM (Vision Language Model)**, can take both Vision and language inputs
- **VLA (Vision Language Action)**, combine vision and language to produce physical actions, often applied in robotics

Use Case Based

- **Gaming**, NPCs in games
- **Healthcare**, Diagnostic agents
- **Manufacturing**, VLA agents for assembly lines
- **Cybersecurity**, Adversarial testing, red teaming
- **Customer Support**, LLM-based agents
- **Software Development**, multi-agent frameworks for writing code

GenAI Agent Interaction Types

NUNNARI LABS

Conversational Agent

- Q&A, Chit Chat, World Knowledge Interactions with Humans
- User query triggered
- Fulfill user queries or transactions

Workflow Agent

- Limited or No Human Interaction
- Event driven triggers
- Fulfill queued tasks or chains of tasks

Workflow Automation -
It’s been there for a decade!

NUNNARI LABS

Traditional Automation

- Workflow automation has been a cornerstone of efficiency and tools like these have led the way in **Robotic Process Automation (RPA)**.
- Example: Automating tasks like **reading an email**, **creating a Jira ticket**.

Workflow Automation - It’s been there for a decade!

NUNNARI LABS

The Next Frontier

- Now, we're entering a new era where these processes are enhanced with **Large Language Models (LLMs)**.
- **Key Innovations :**
  - **Natural Language Processing (NLP):** Automating complex workflows with minimal coding.
  - **Prompting & Few-Shot Learning:** Rapidly deploying automation using examples rather than extensive programming.

NUNNARI
LABS

Foundational Components of Agents!

Navaneeth Malingan

Primary Components

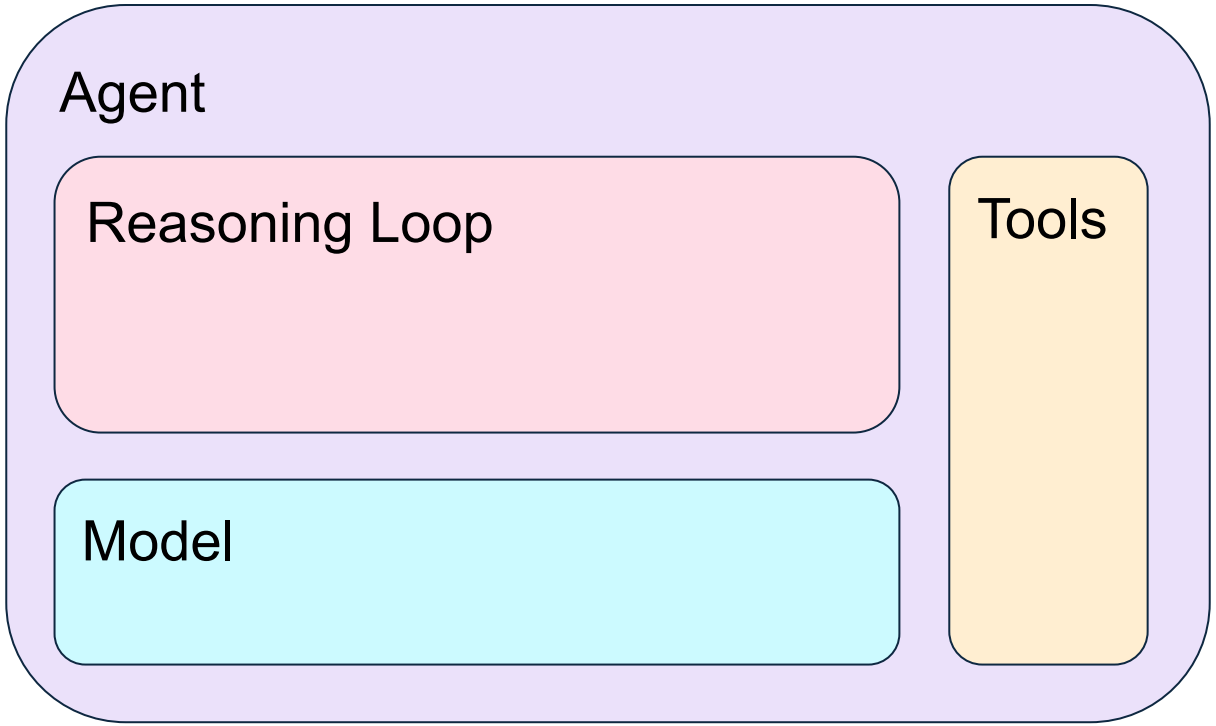

Model

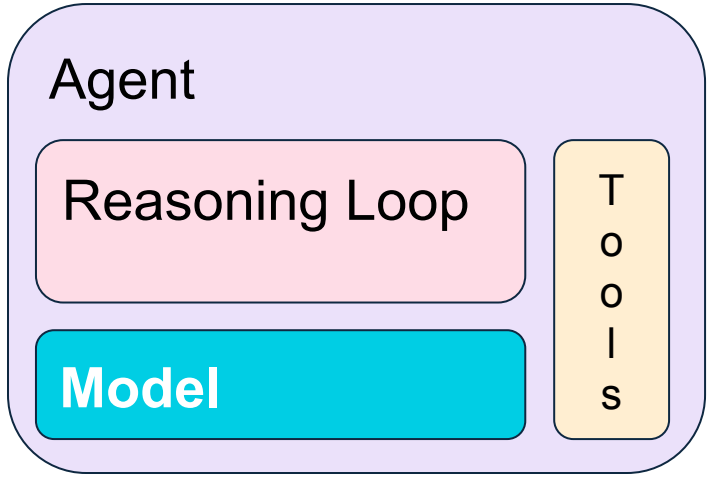

- **Any Generative Language model (large or small)**

  `gpt-4, gemini-pro, claude, etc.`

- **Models can be multimodal or fine-tuned**

  `gemini-pro-vision, DALL-E, etc.`

- **Training Data Can Play a Large Role**

  Preference for Models trained with data signatures from Tools and Reasoning steps

Tools

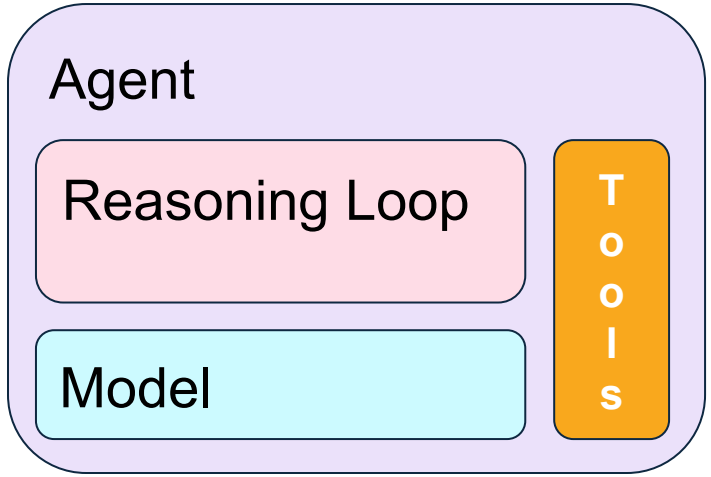

- Allow Agents to interact with **external data and services**
- 3 primary types
  - Extensions • Functions • Data Stores

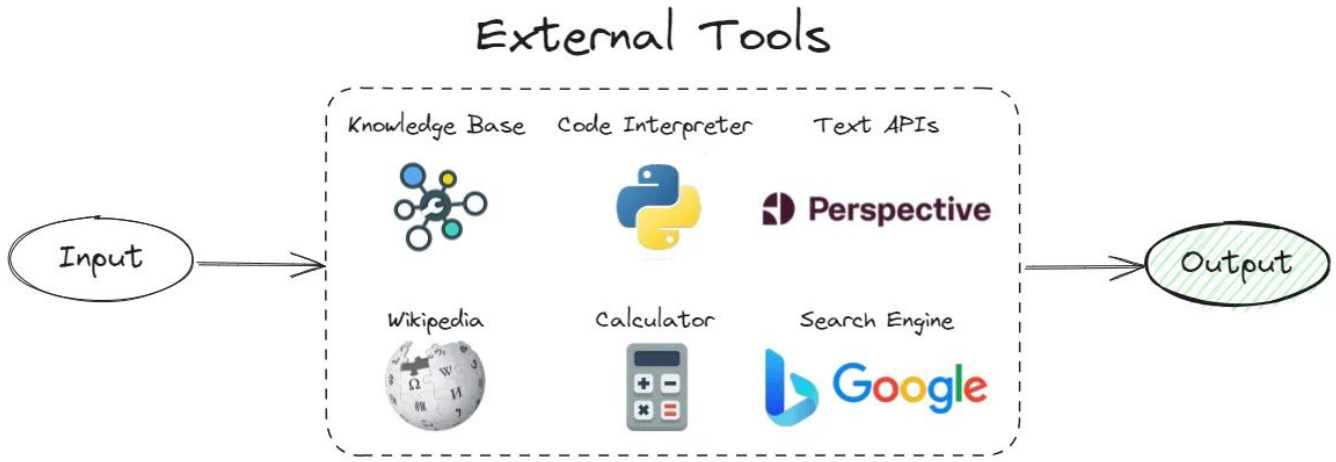

Reasoning Loop

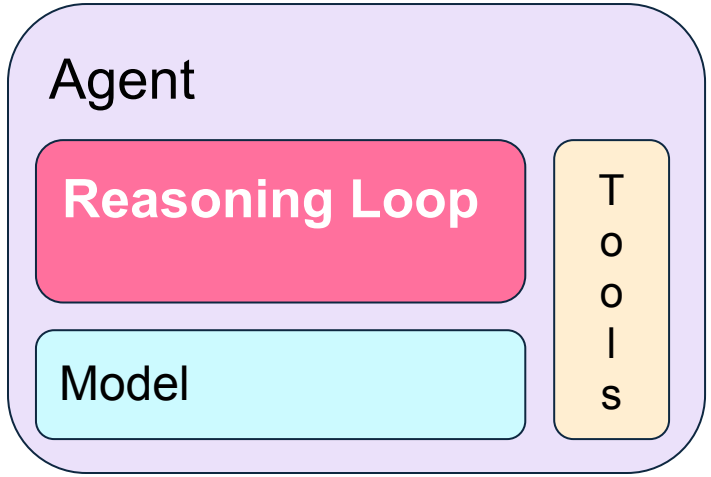

- **Iterative, Introspective process**

  Aimed at taking actions towards achieving a goal
- **Advanced Prompt Engineering Frameworks**
  - *Simple rule-based calculations*
  - *Complex thought chains*
  - *ML algorithms*
  - *Probabilistic reasoning techniques*

- **Chain-of-Thought**
- **Tree-of-Thoughts**

- **Directional Stimulus Prompting**
- **ReAct Agent**

Chain of Thought Prompting

NUNNARI LABS

- CoT prompting is a simple technique for improving an LLM's performance on reasoning tasks like commonsense or symbolic reasoning.
- It leverages few-shot learning by inserting several examples of reasoning problems being solved within the prompt.

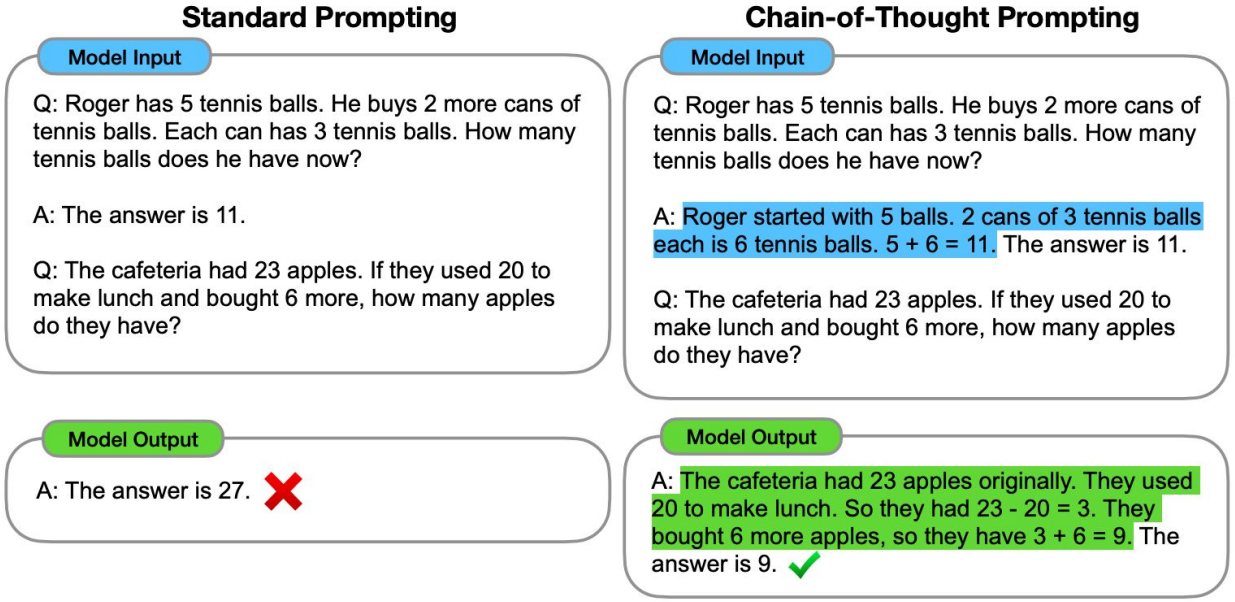

Figure 1: Chain-of-thought prompting enables large language models to tackle complex arithmetic, commonsense, and symbolic reasoning tasks. Chain-of-thought reasoning processes are highlighted.

Tree of Thoughts Prompting

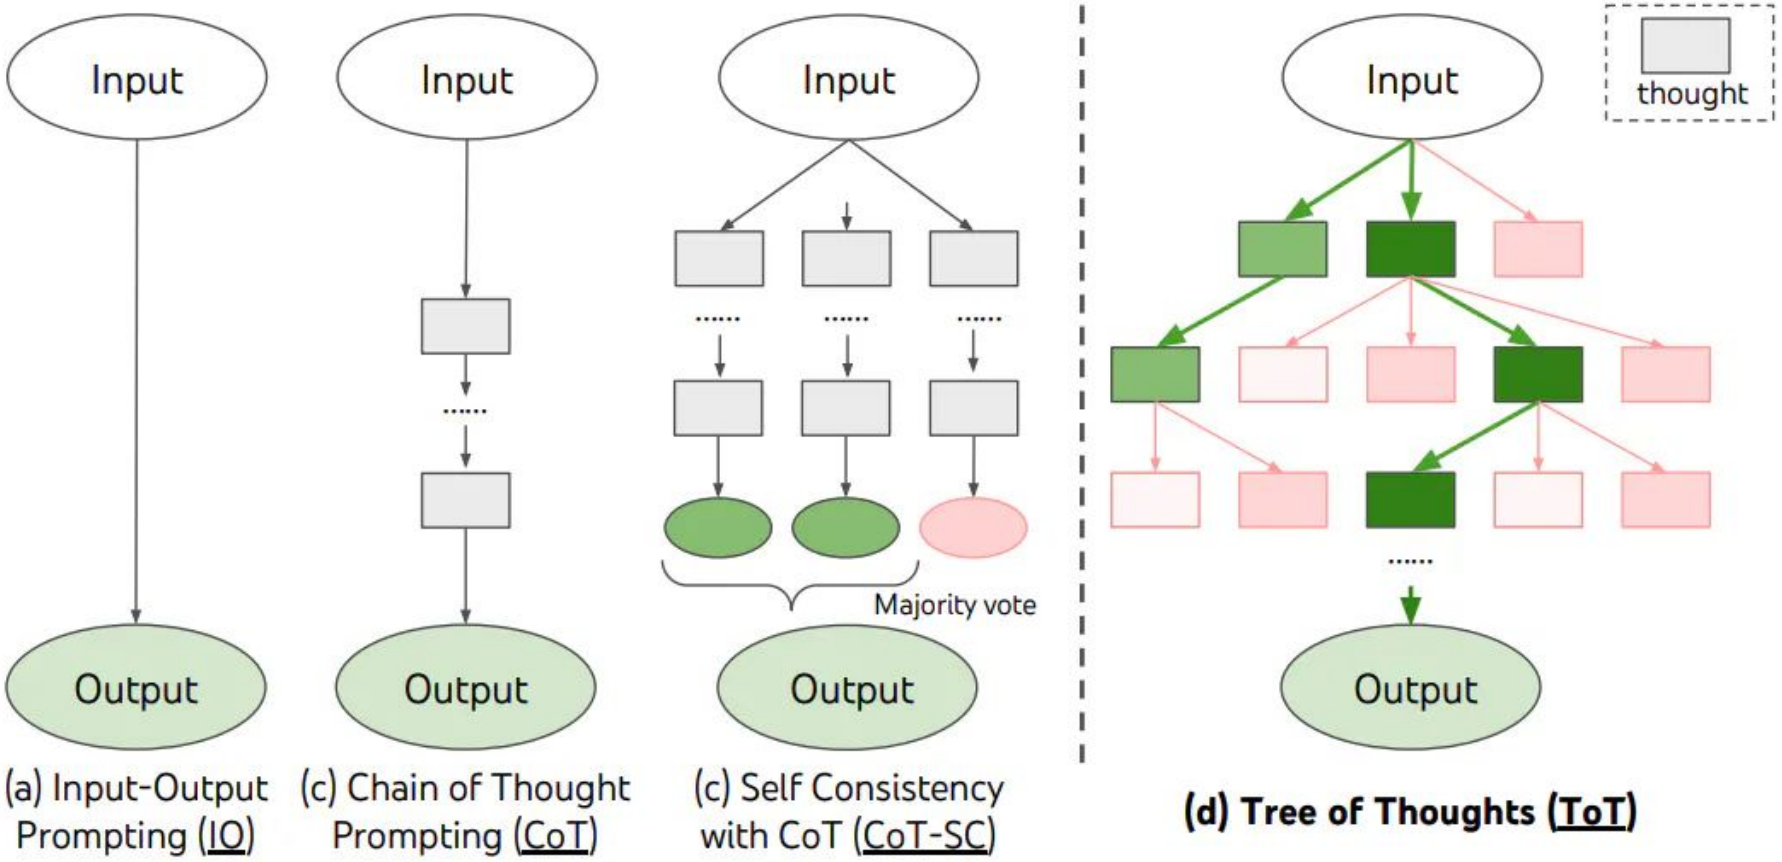

Directional Stimulus Prompting

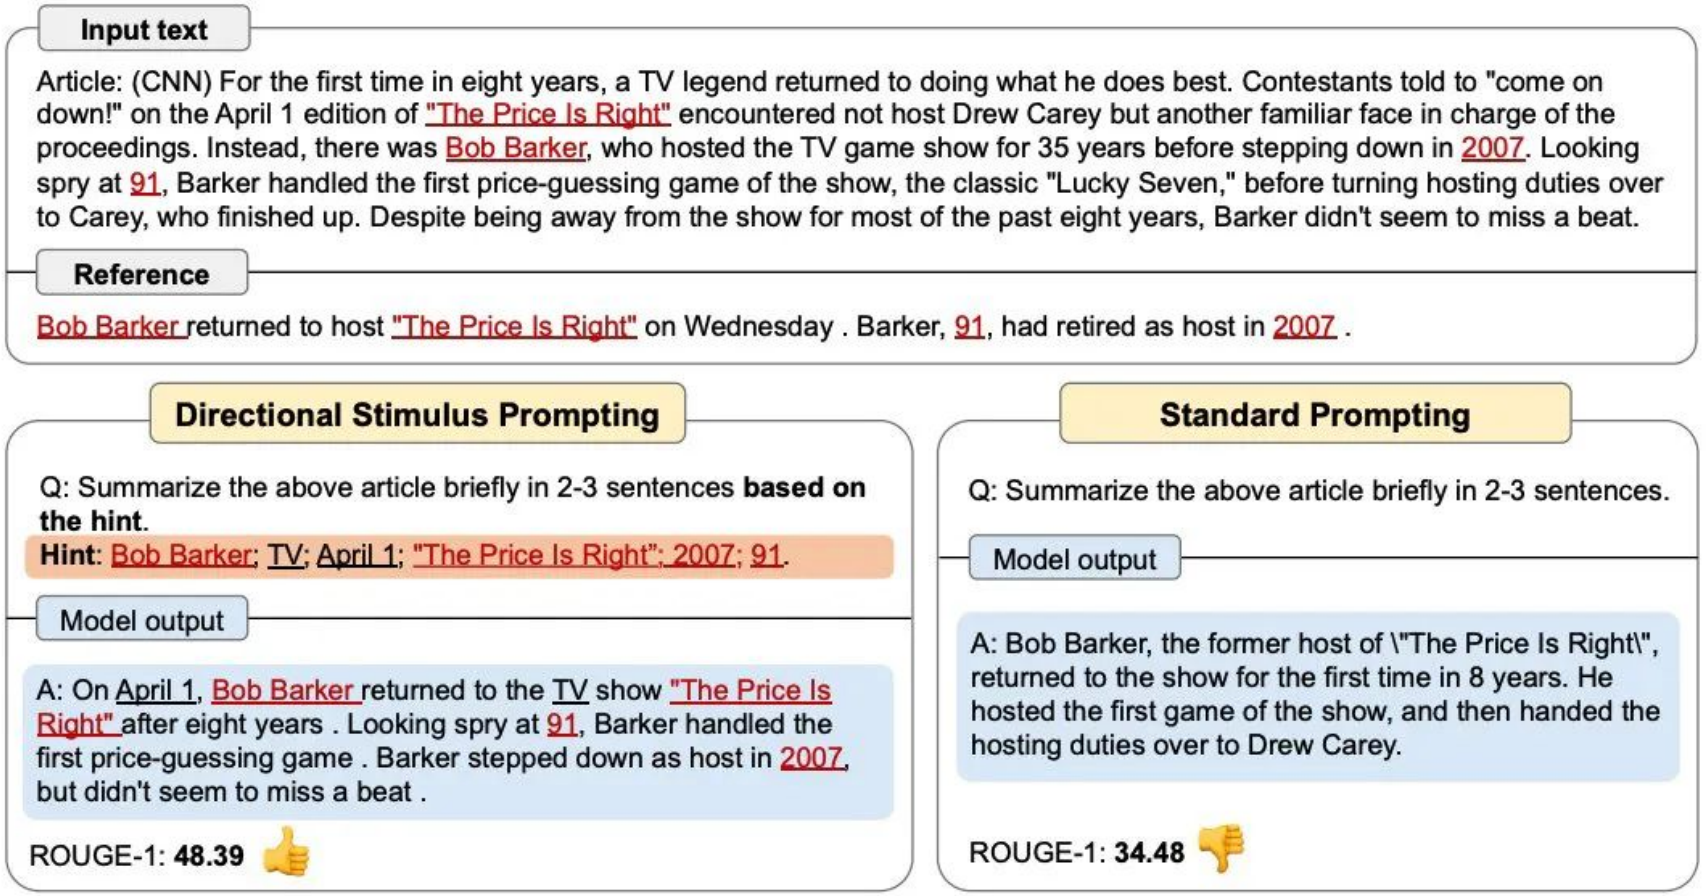

ReAct Agent

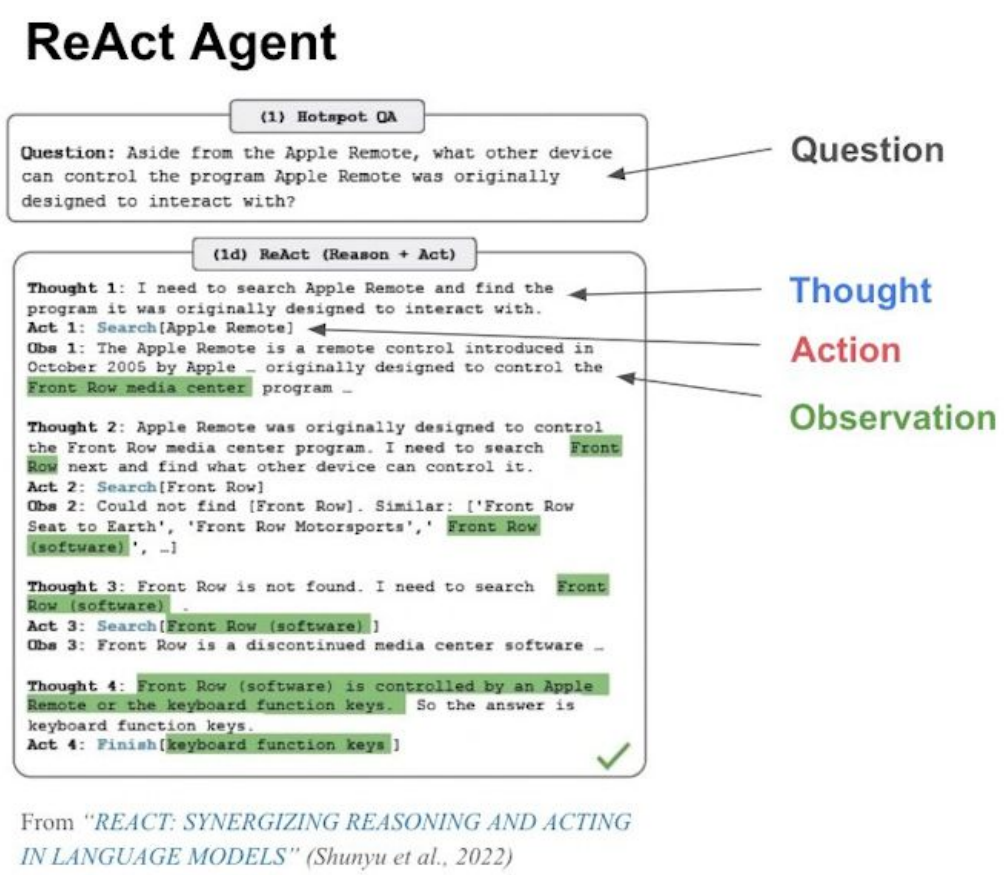

From *"REACT: SYNERGIZING REASONING AND ACTING IN LANGUAGE MODELS" (Shunyu et al., 2022)*

```
PROMPT_TEMPLATE = <<~PROMPT
  Today is %{today} and you can use tools to get new information. Answer the question as best as you can using the following tools:

  %{tool_description}

  Use the following format:

  Question: the input question you must answer
  Thought: comment on what you want to do next
  Action: the action to take, exactly one element of [%{tool_names}]
  Action Input: the input to the action
  Observation: the result of the action
  ... (this Thought/Action/Action Input/Observation repeats N times, use it until you are sure of the answer)
  Thought: I now know the final answer
  Final Answer: your final answer to the original input question

  Begin!

  Question: %{question}
  Thought: %{previous_responses}
PROMPT
```

ReAct Agent in 150 lines of code

NUNNARI
LABS

Models vs Agents

Models	Agents
Knowledge is limited to what is available in their training data.	Knowledge is extended through the connection with external systems via Tools
OOTB single turn inference / prediction based on the user query.	OOTB multi turn inference / prediction based on decisions made in the reasoning loop.
No native tool implementation. Tools can be implemented via custom integrations.	Tools are natively implemented in Agent architecture.
No OOTB logic layer implemented. Users can form prompts as simple questions or use reasoning frameworks (CoT, ReAct) to form complex prompts to guide the model in prediction.	OOTB cognitive architecture that use reasoning frameworks like CoT, ReAct, or other pre-built Agent frameworks like LangChain.

NUNNARI
LABS

Agent Basic Operation

Navaneeth Malingan

User Queries or Events Initiate Interactions

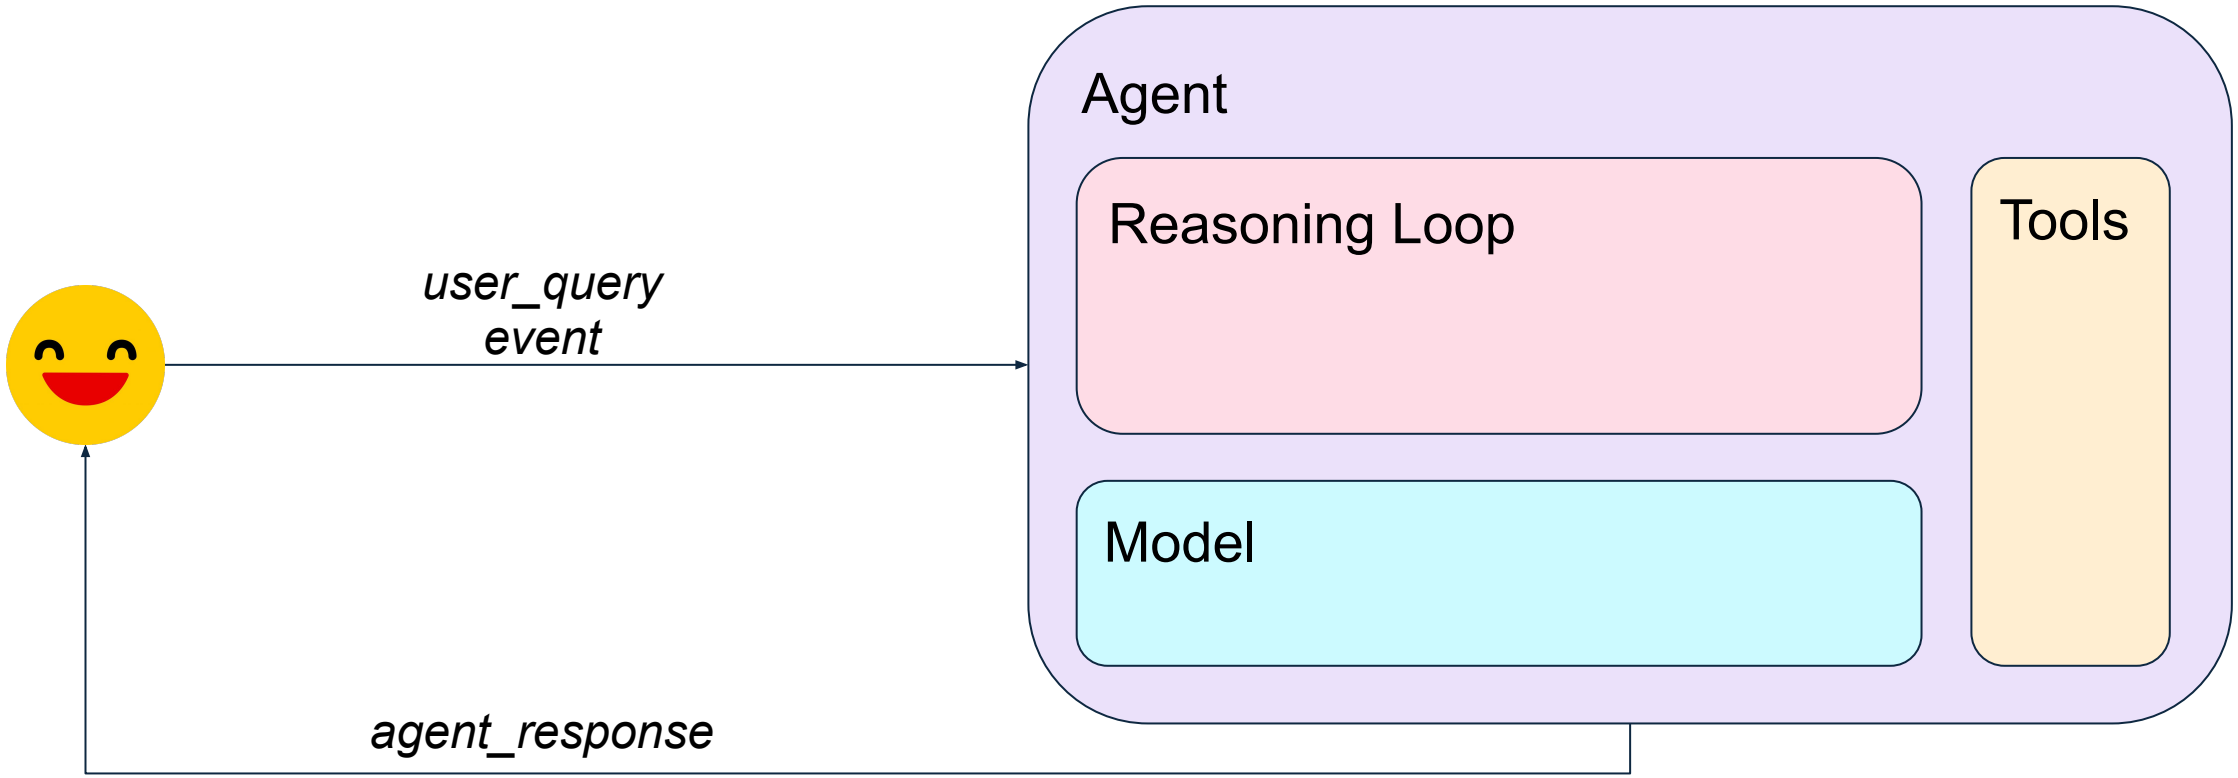

NUNNARI LABS

User Queries or Events Initiate Interactions

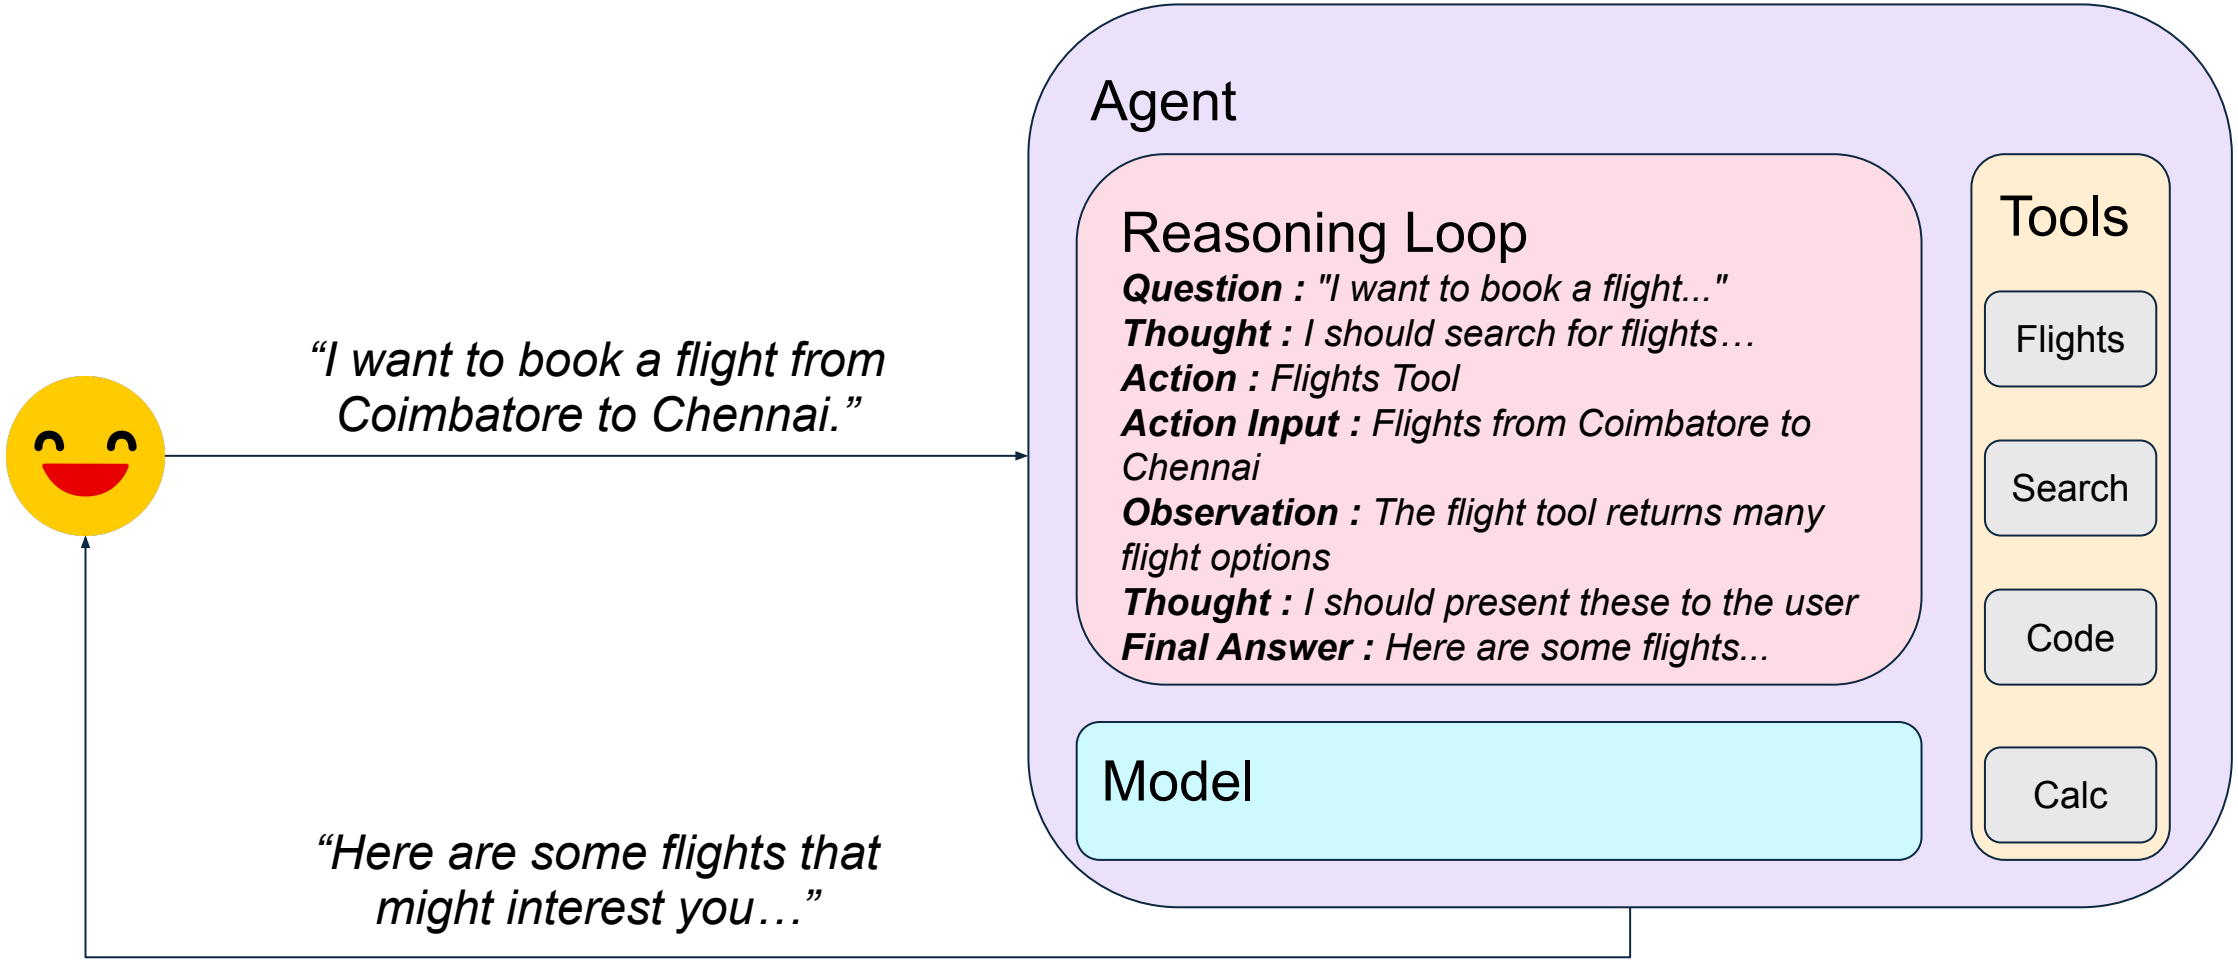

Tools : Extensions

Navaneeth Malingan

Connecting Agents to APIs

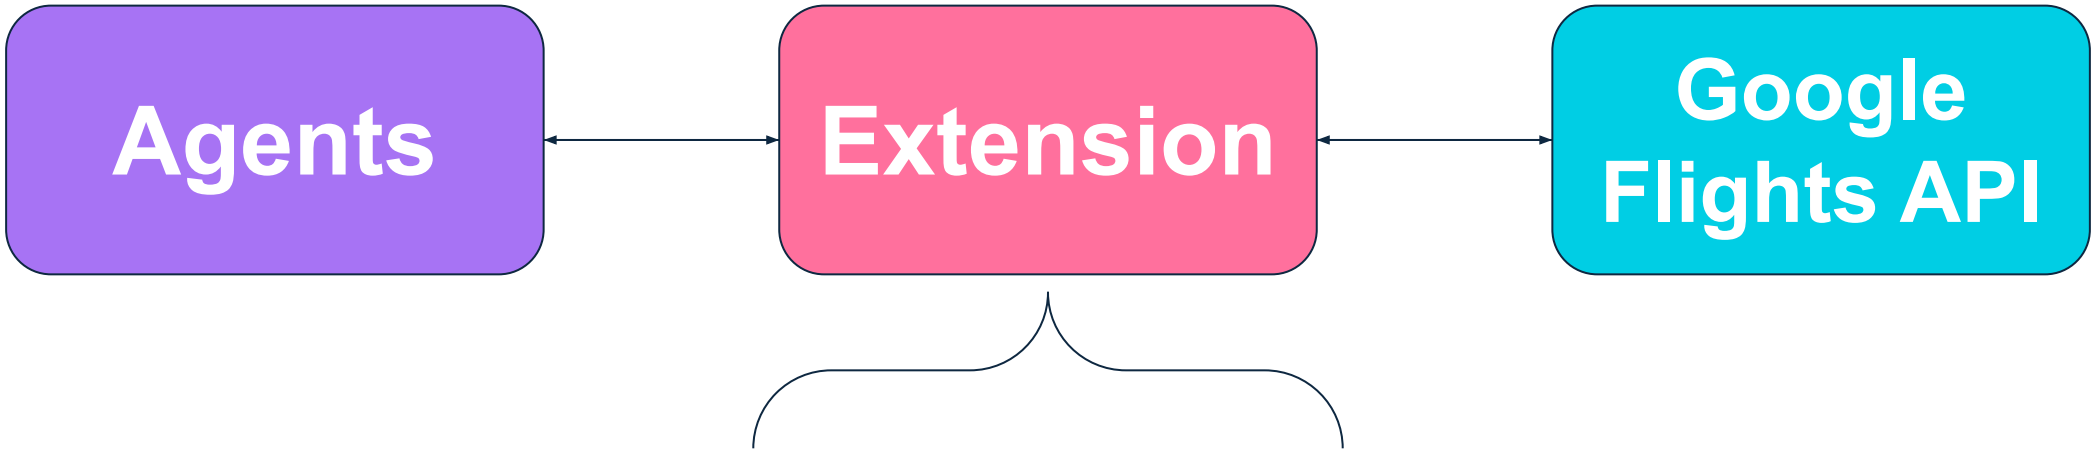

1. "When the user wants to search for flights, call `get_flights`..." (WHEN)
2. "Input args for `get_flights` are `arg1, arg2,...`" (HOW)
3. "The `get_flights` method can be used to get the latest..." (WHAT TO EXPECT)

NUNNARI
LABS

Connecting Agents to APIs

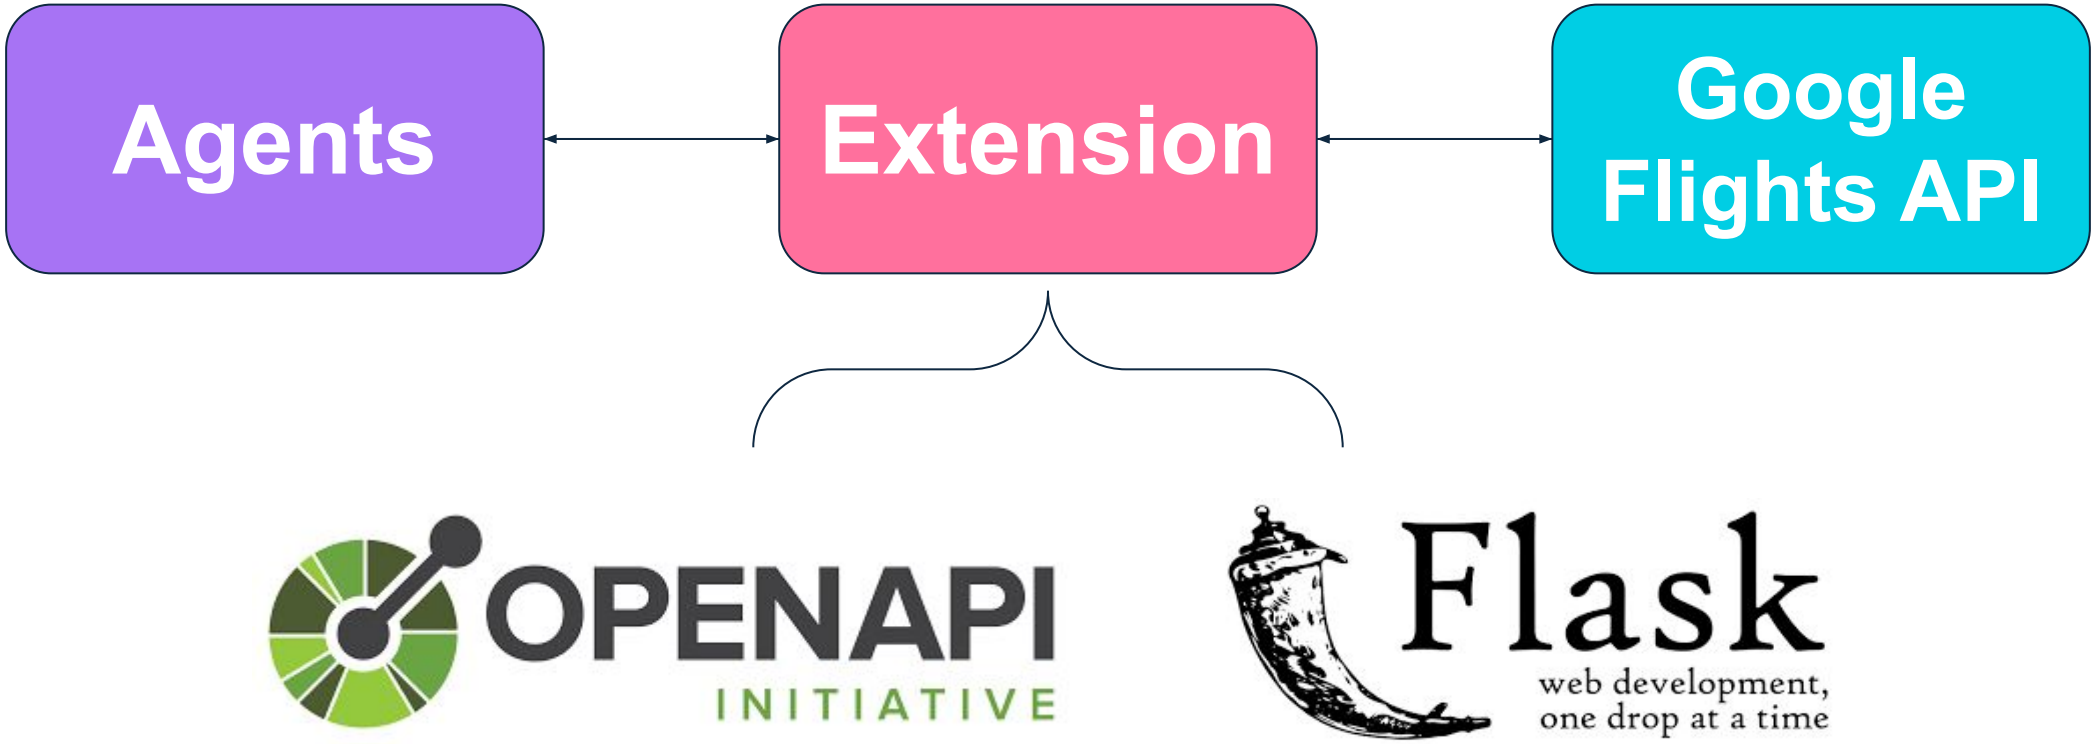

NUNNARI
LABS

Tools : Function Calling

Navaneeth Malingan

Stubbing APIs & Division of Labour

NUNNARI LABS

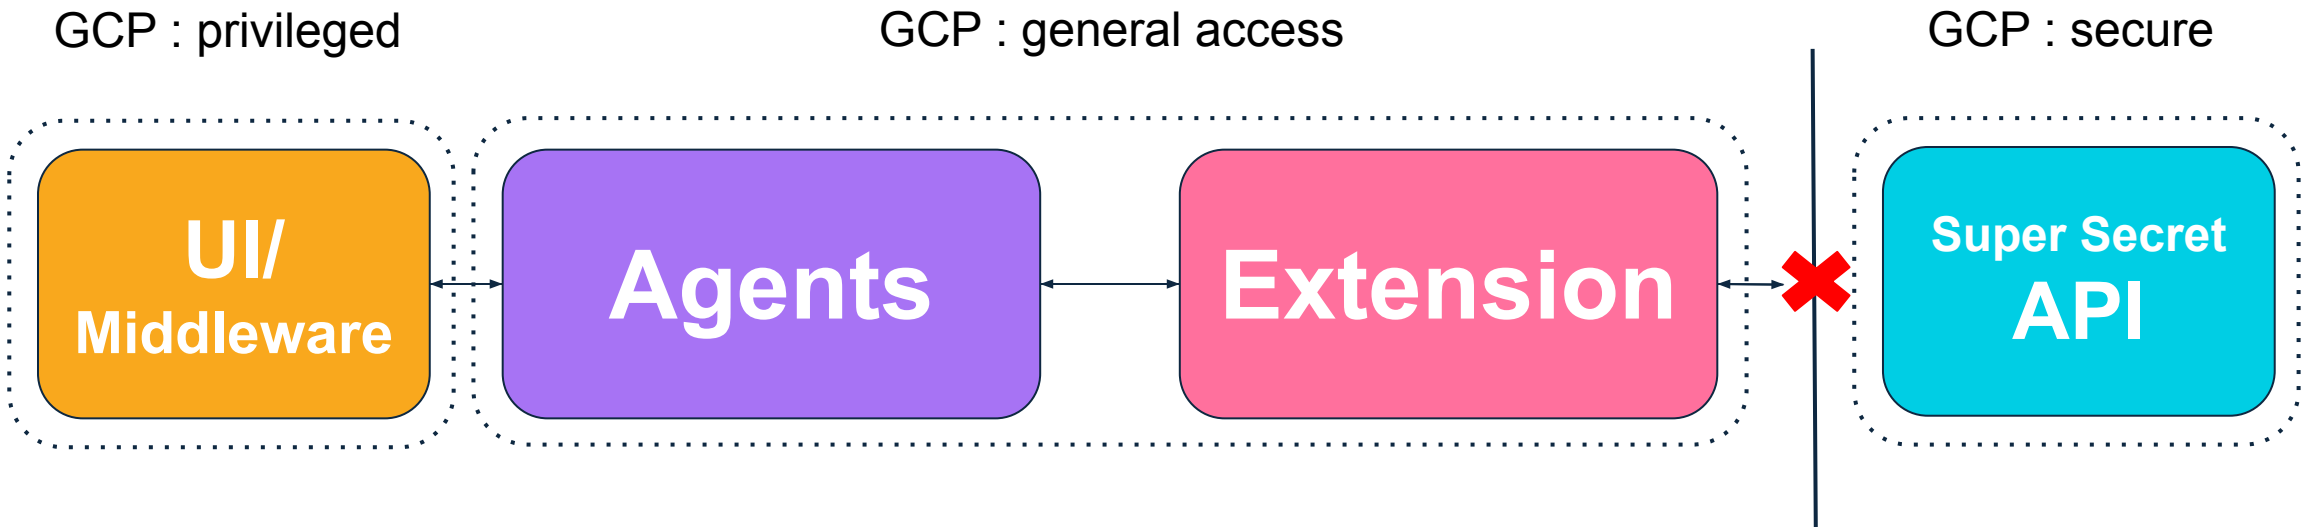

Stubbing APIs & Division of Labour

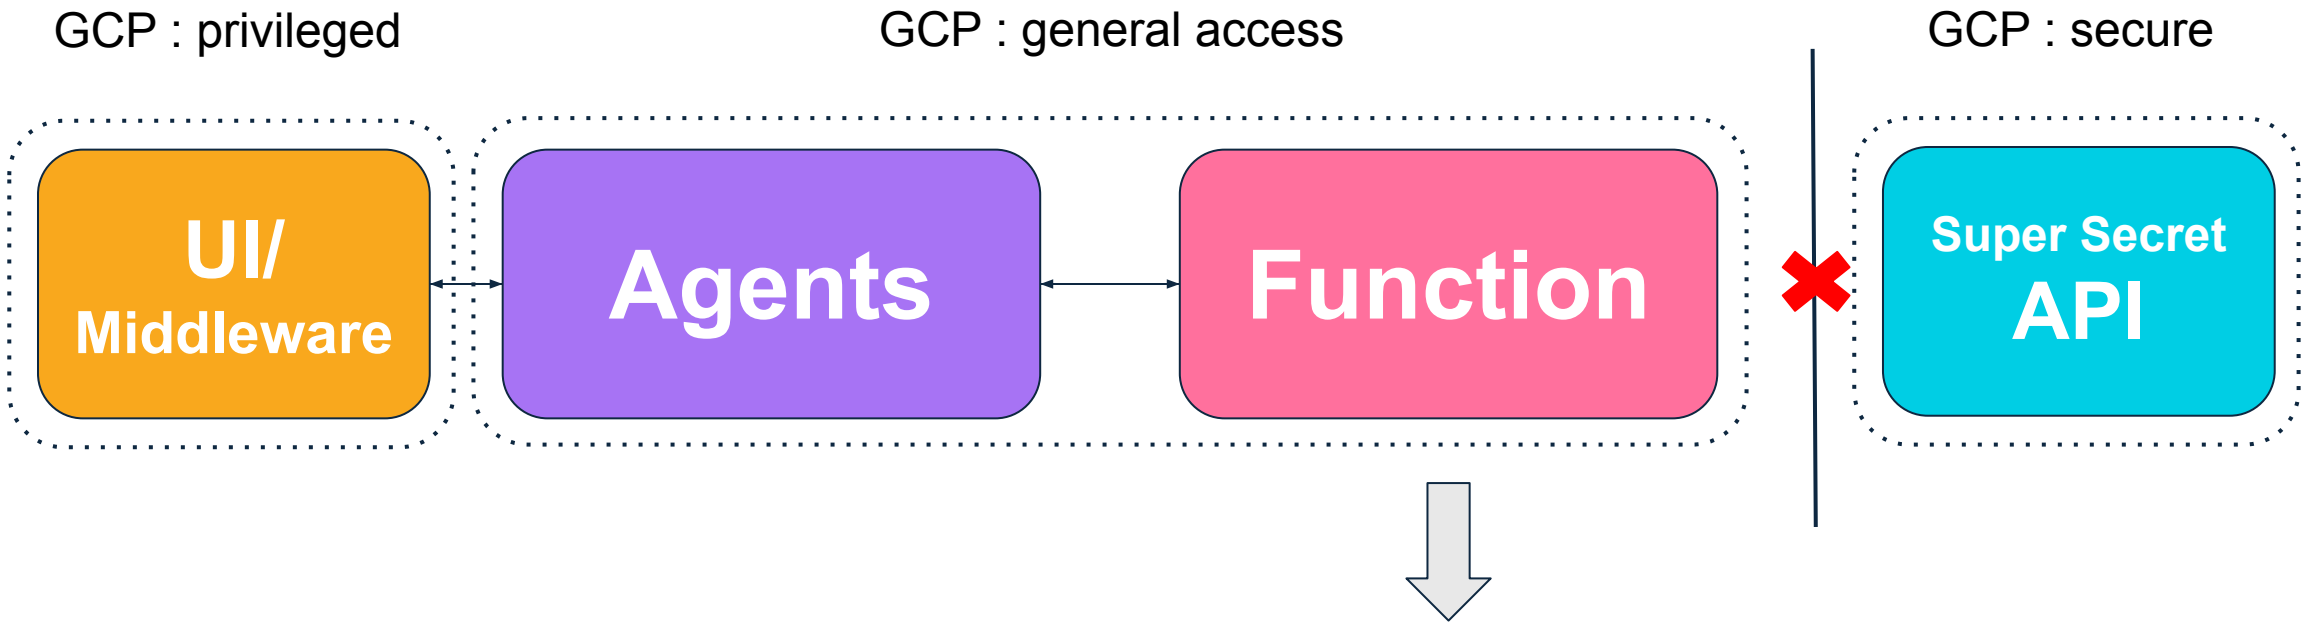

1. "The user wants to search for flights, call `get_flights...`" (WHEN)
2. Return: `{"function_call":{"name": "get_flights"...}}` (HOW)

Stubbing APIs & Division of Labour

NUNNARI LABS

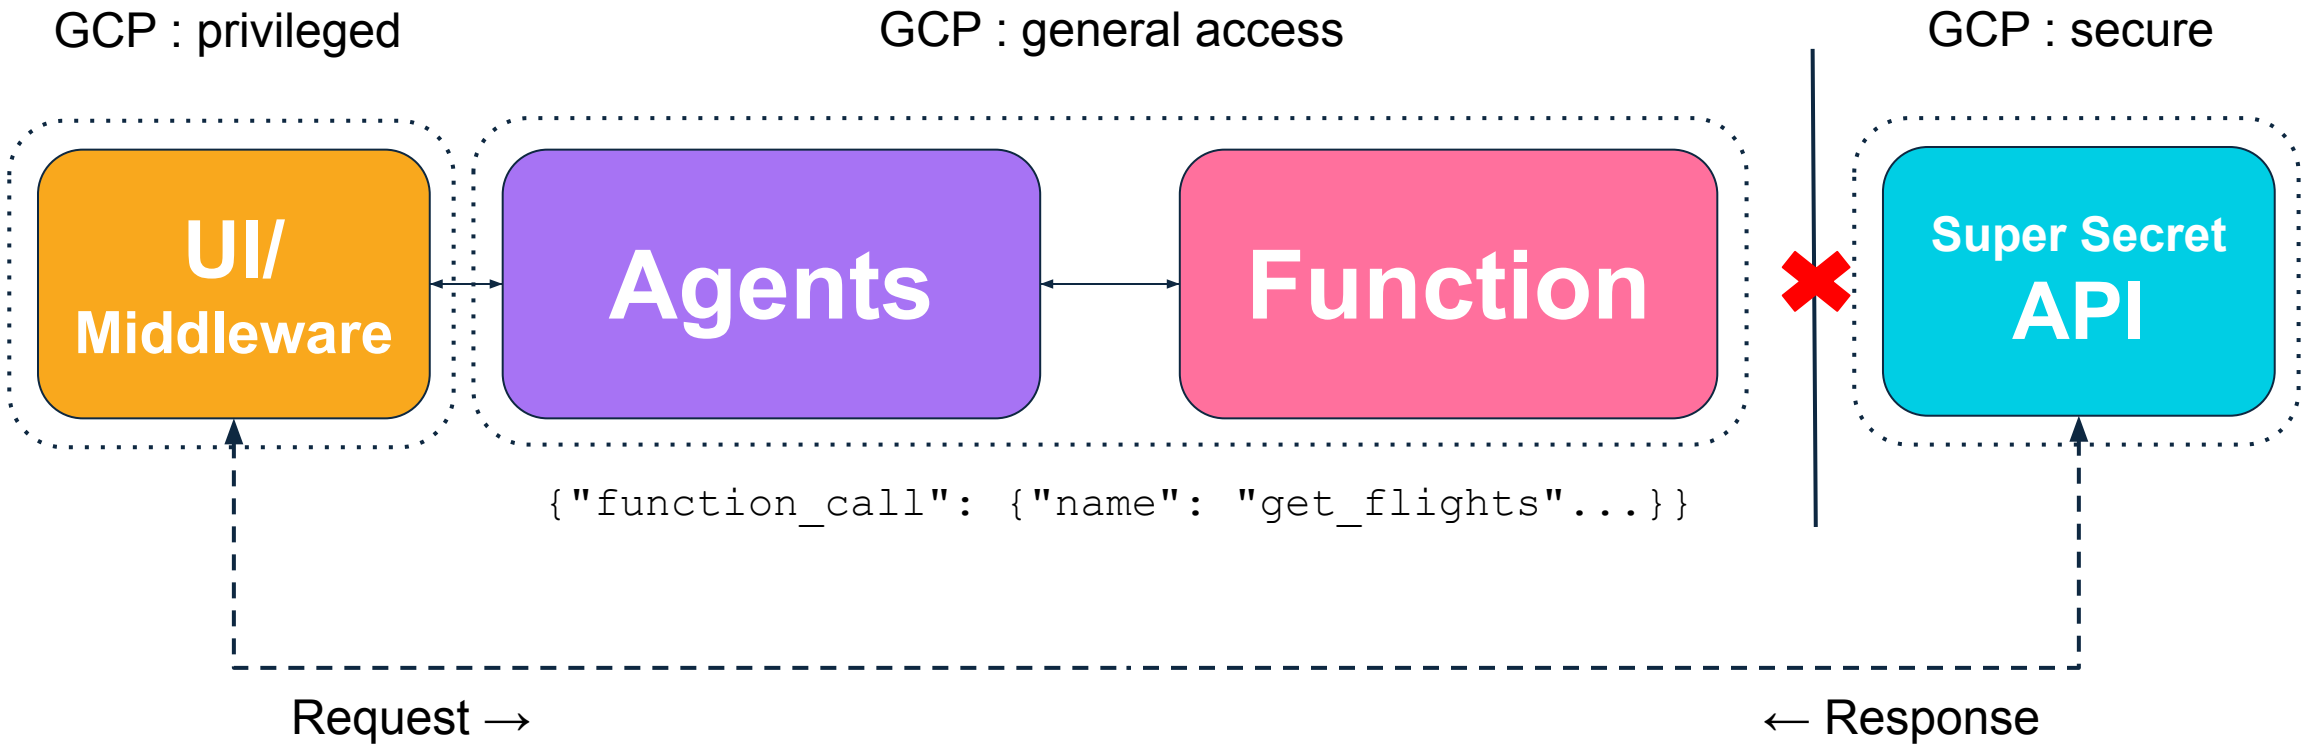

NUNNARI
LABS

Extensions vs Functions

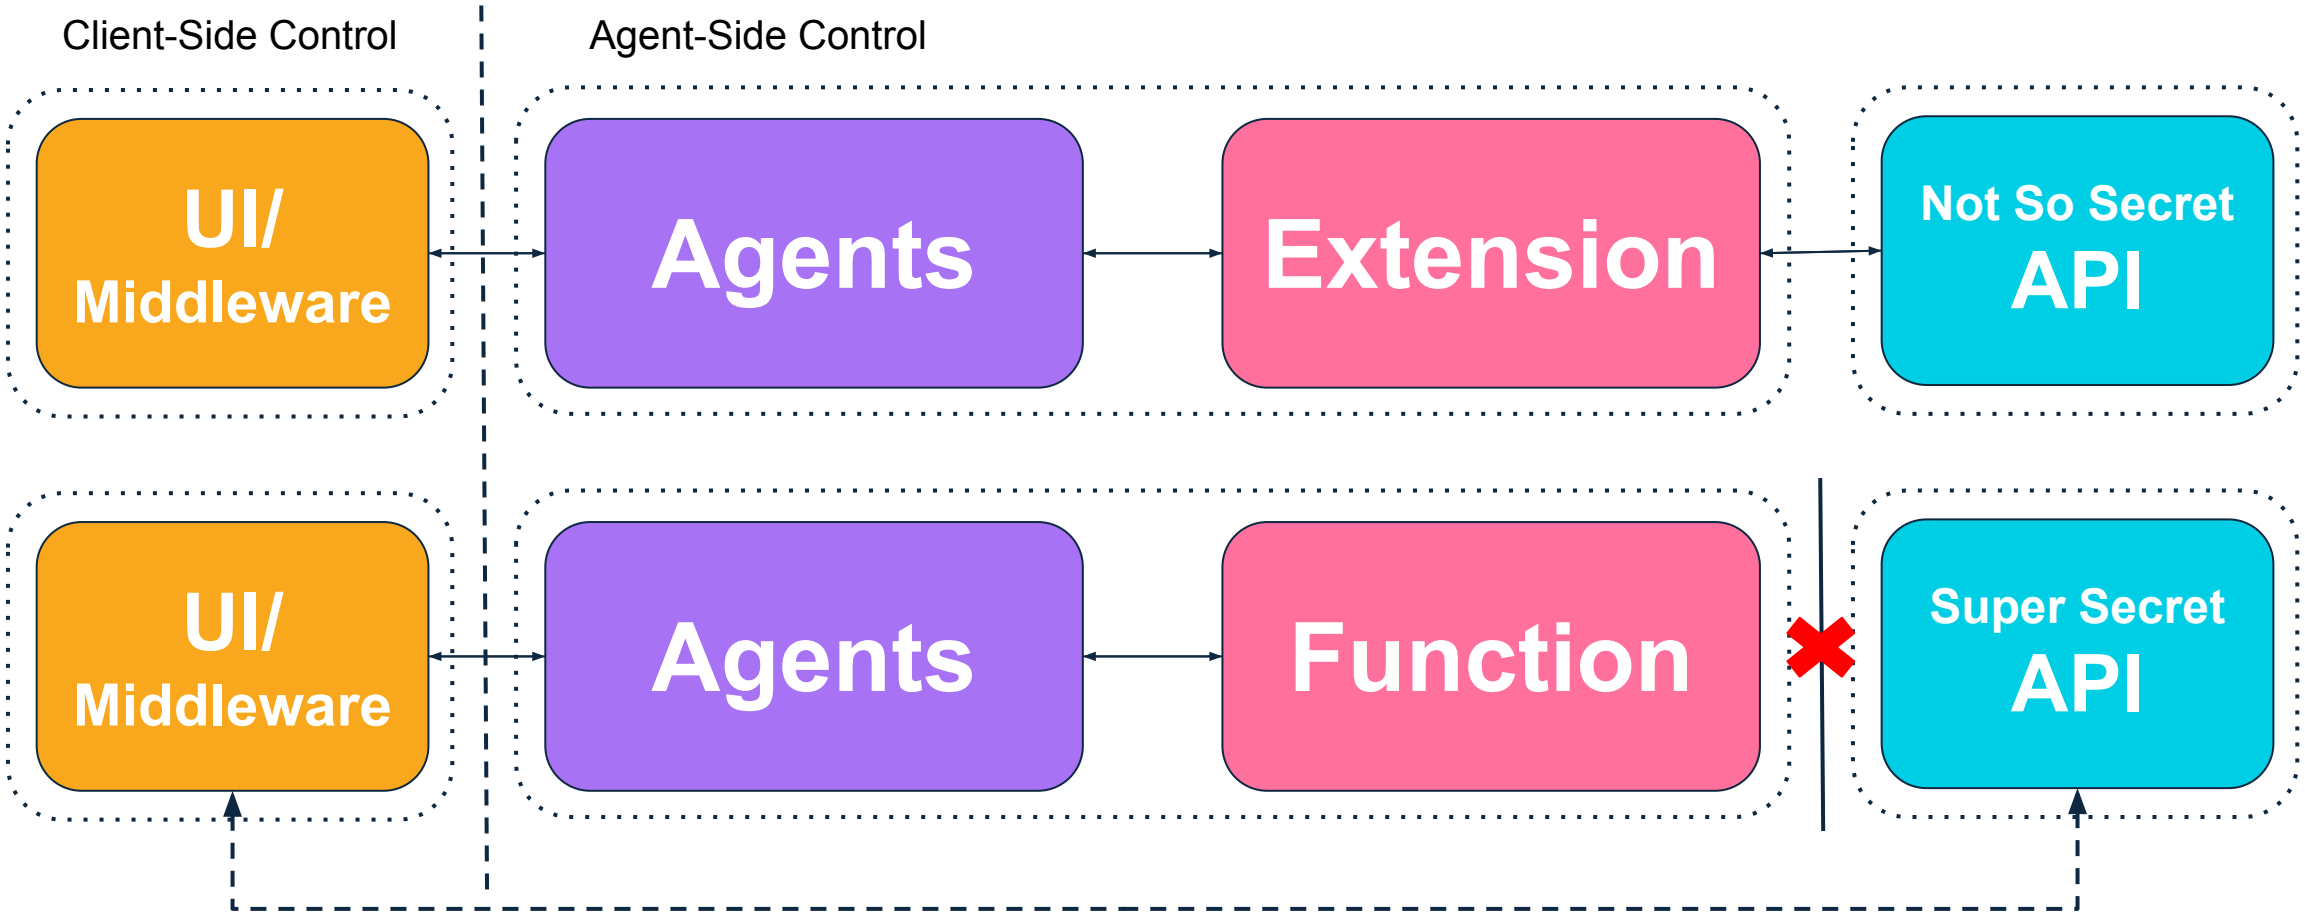

Tools : Data Stores

Navaneeth Malingan

NUNNARI
LABS

Indexing & Vector Databases

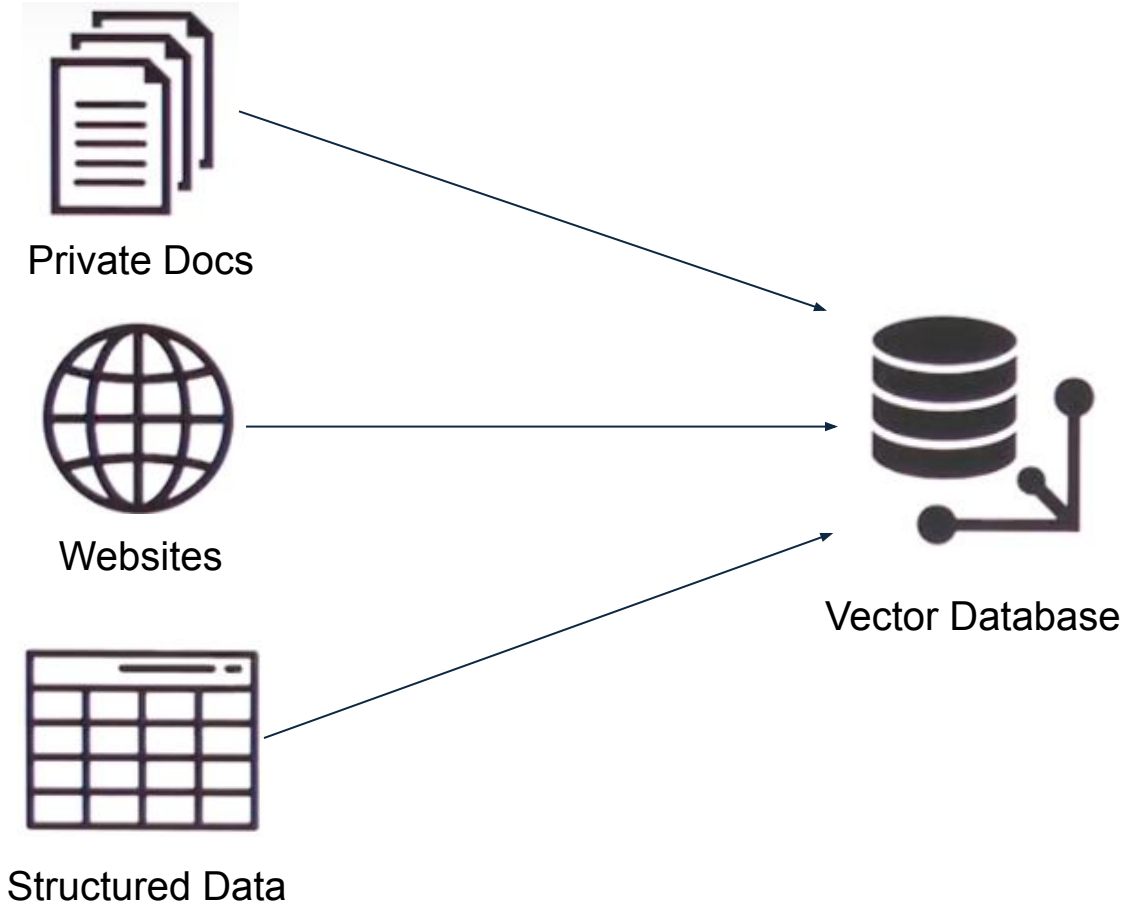

NUNNARI
LABS

Indexing & Vector Databases

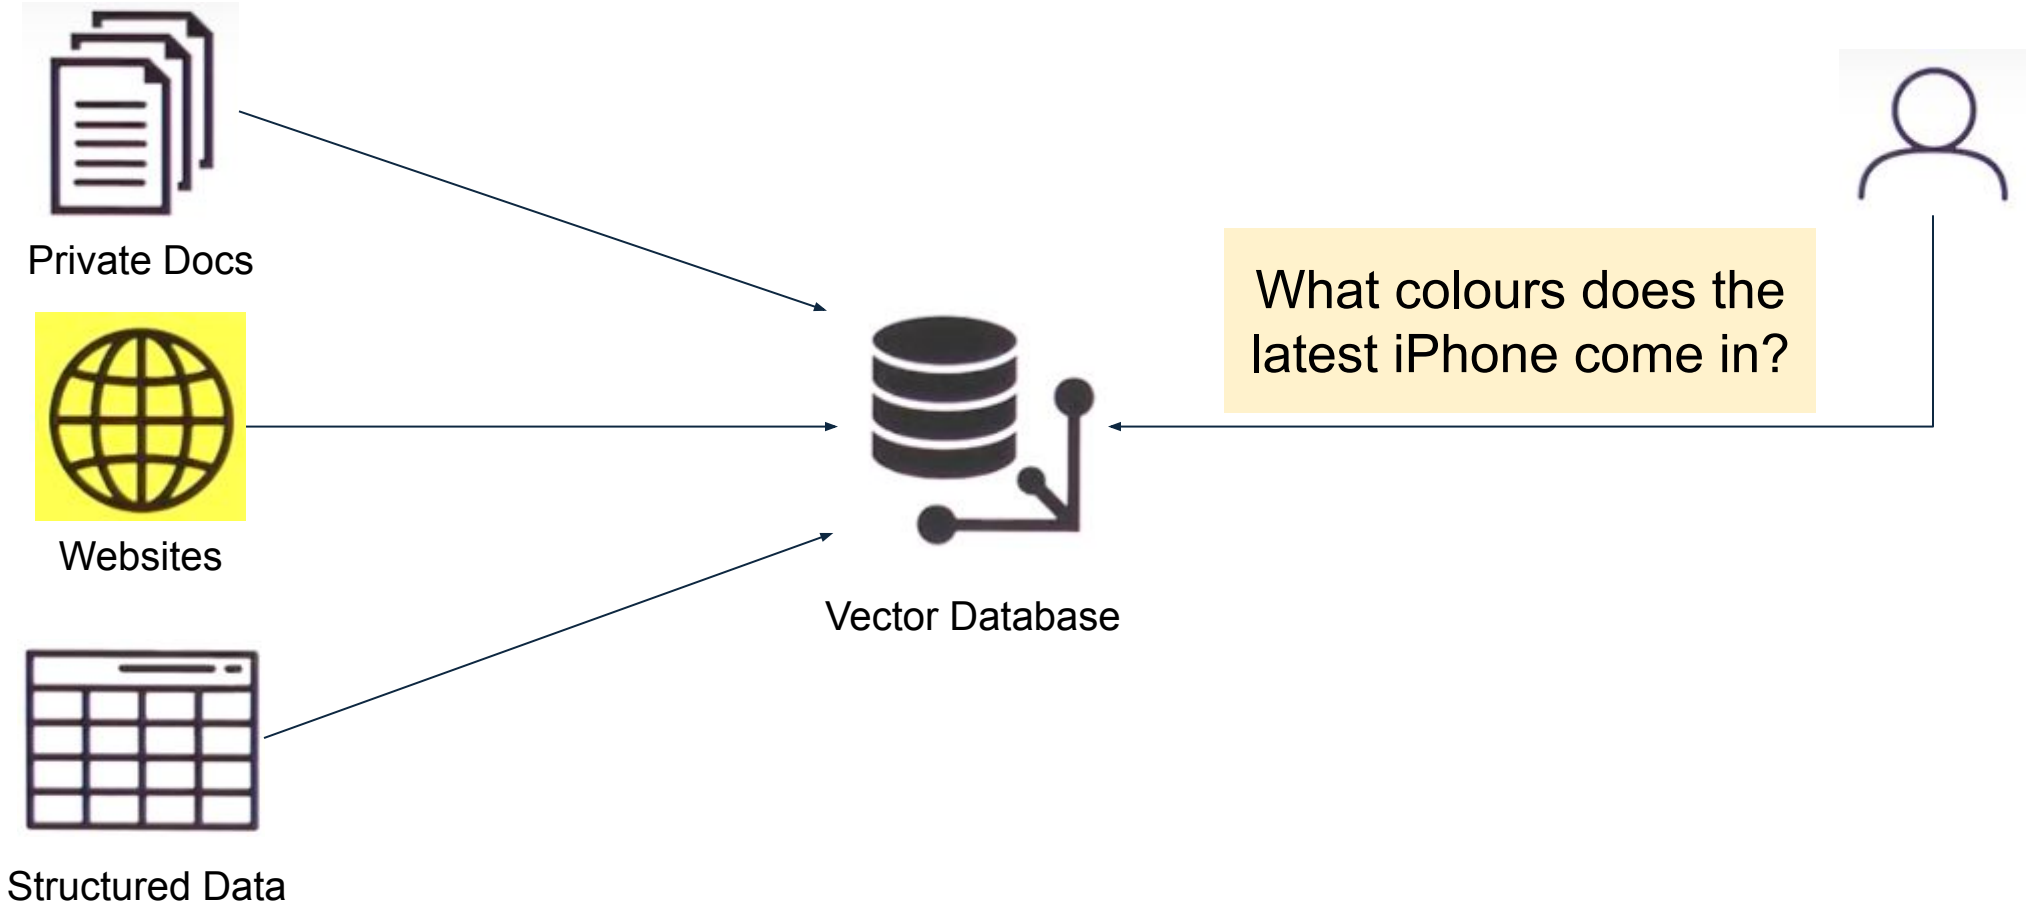

Indexing & Vector Databases

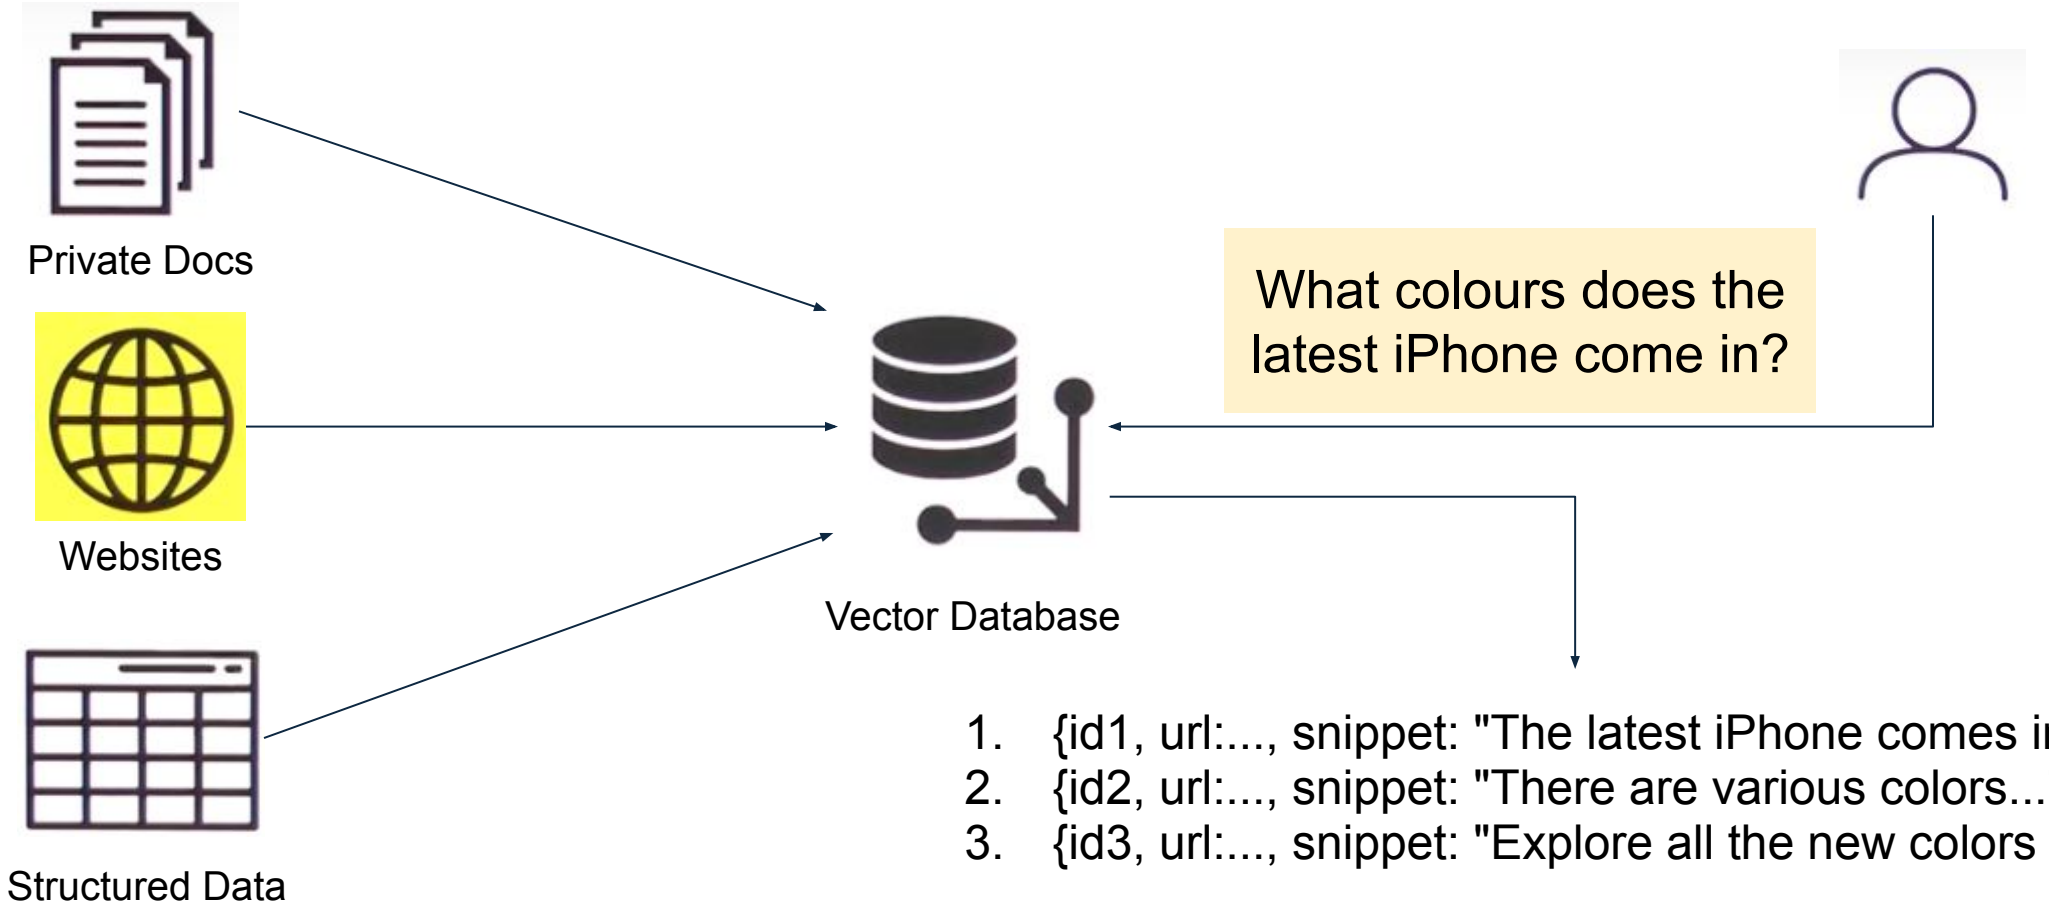

NUNNARI
LABS

Sample Architecture

Navaneeth Malingan

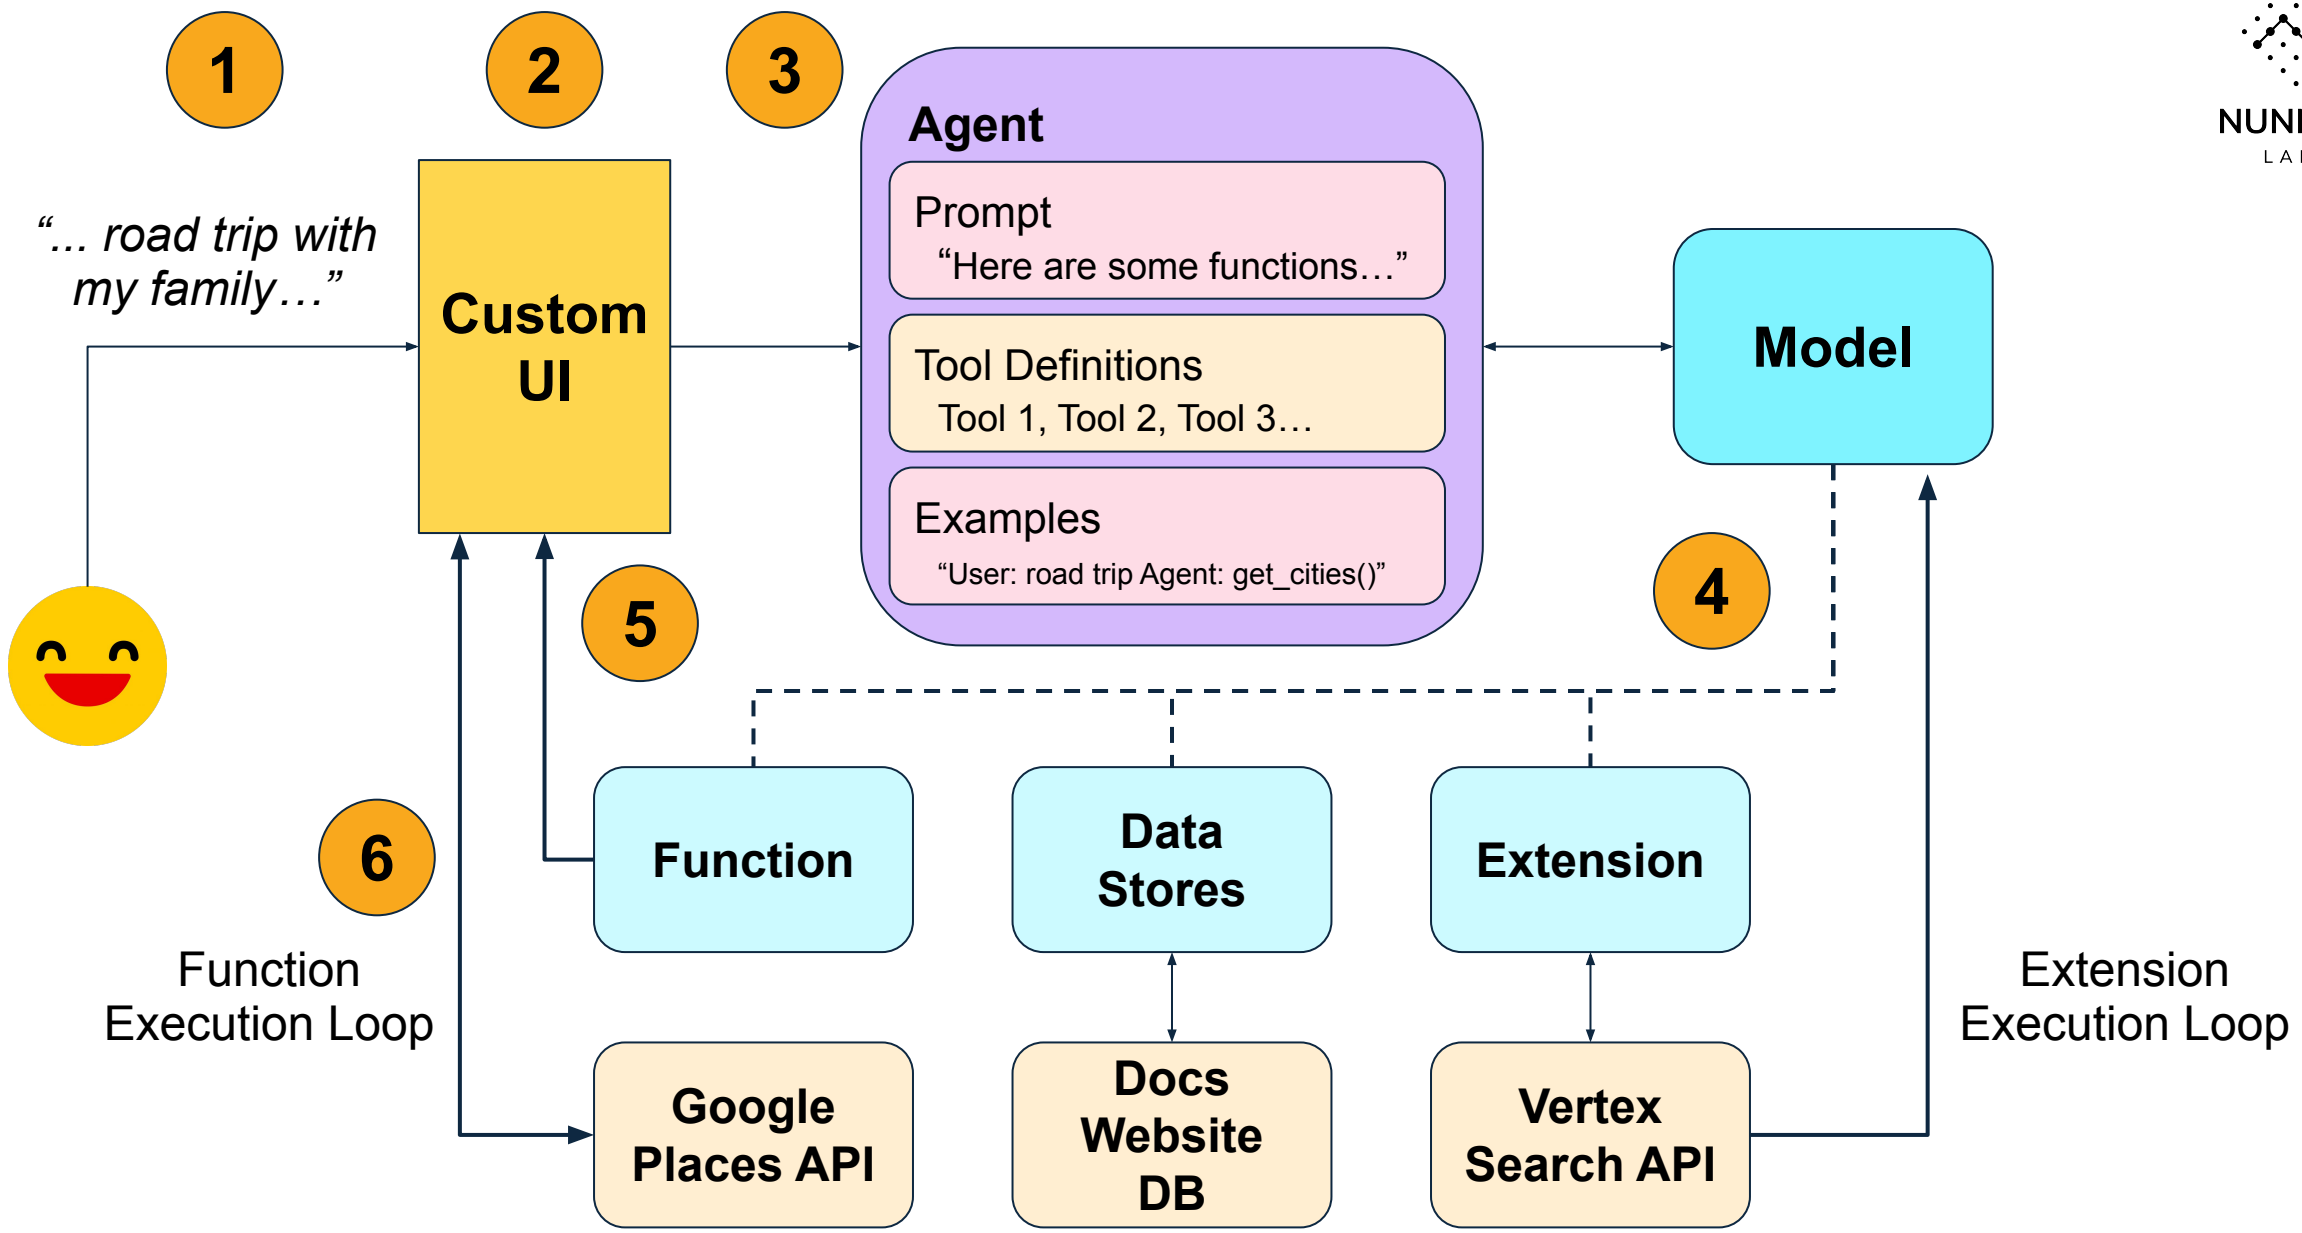

NUNNARI LABS

Sample Architecture

Navaneeth Malingan

NUNNARI
LABS

Let’s Sum up Shortly!

Navaneeth Malingan

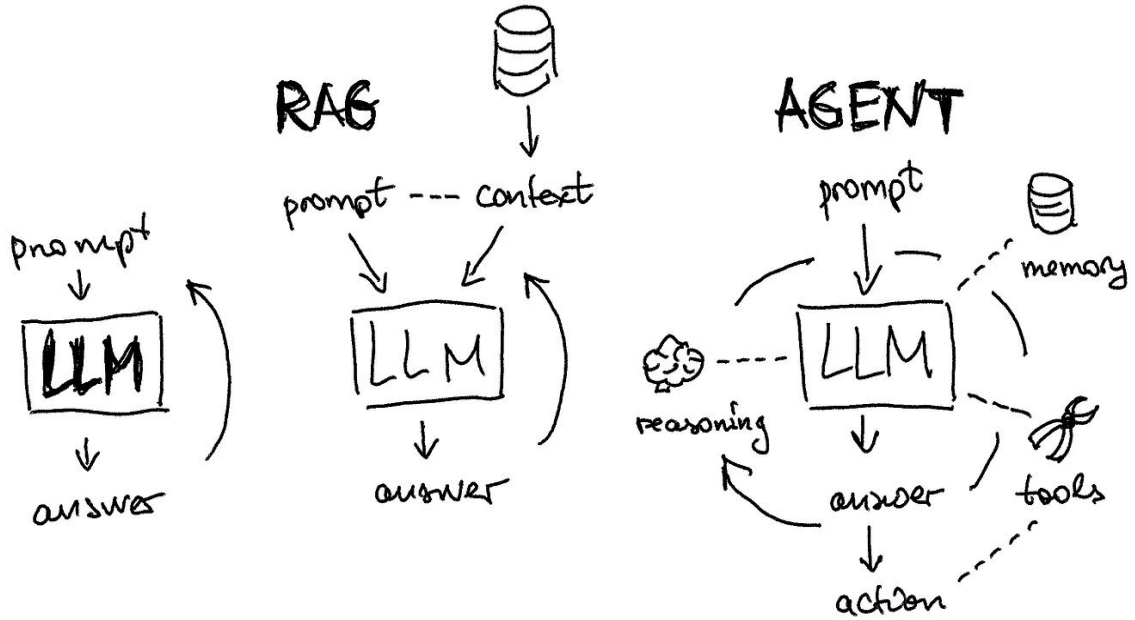

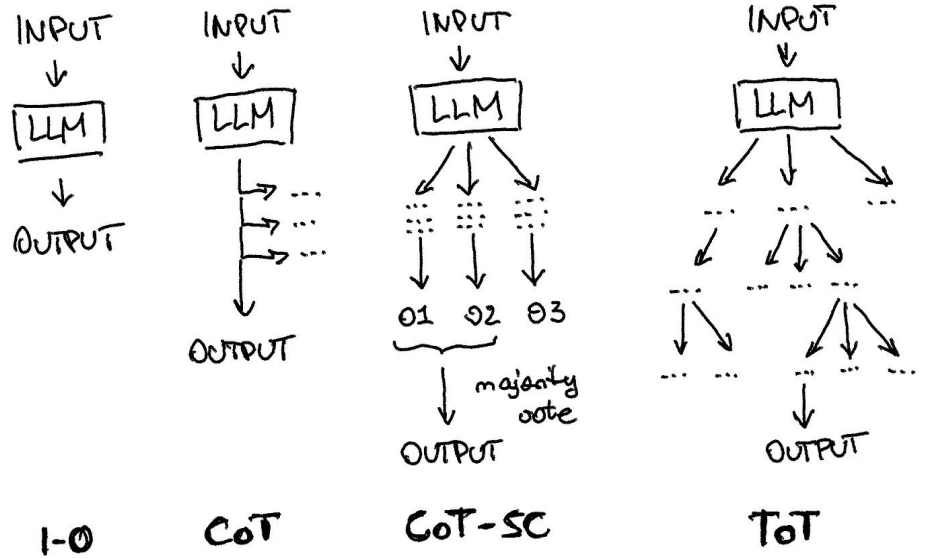

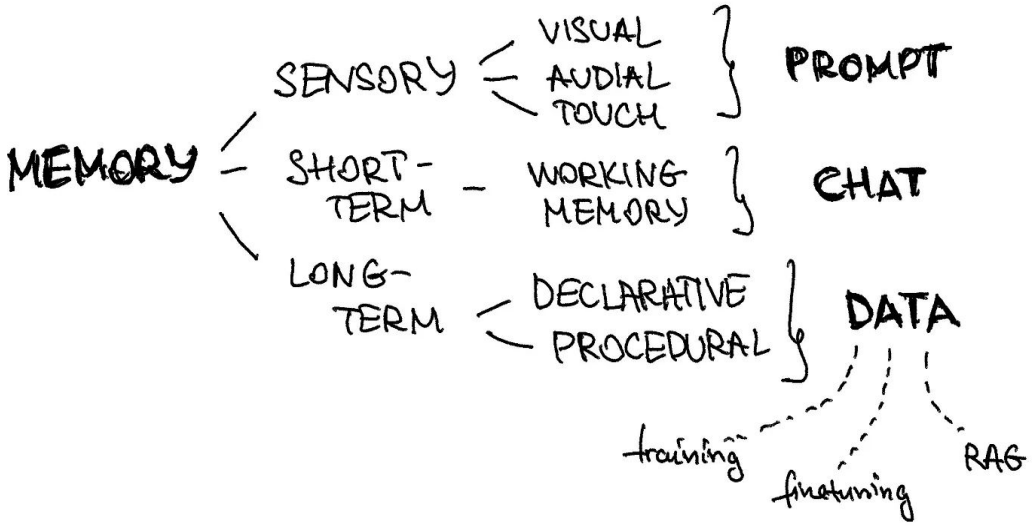

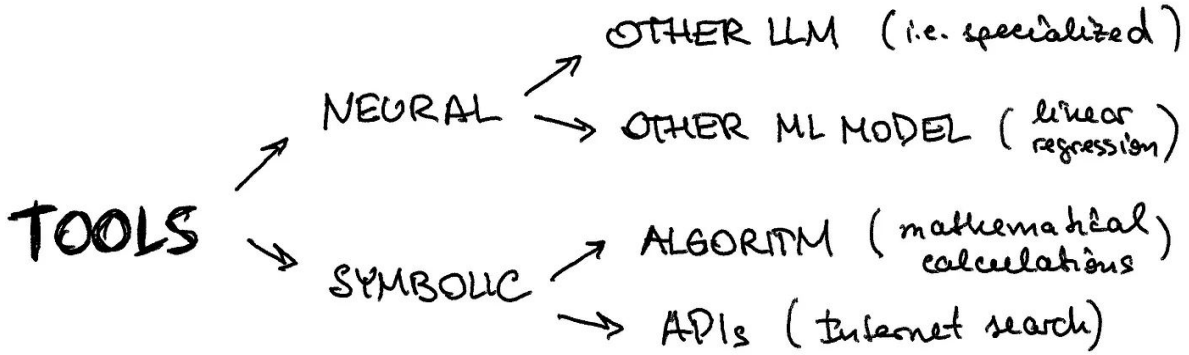

Navaneeth Malingan

NUNNARI
LABS

Multi-Agent Architecture

Navaneeth Malingan

NUNNARI
LABS

Collaborative

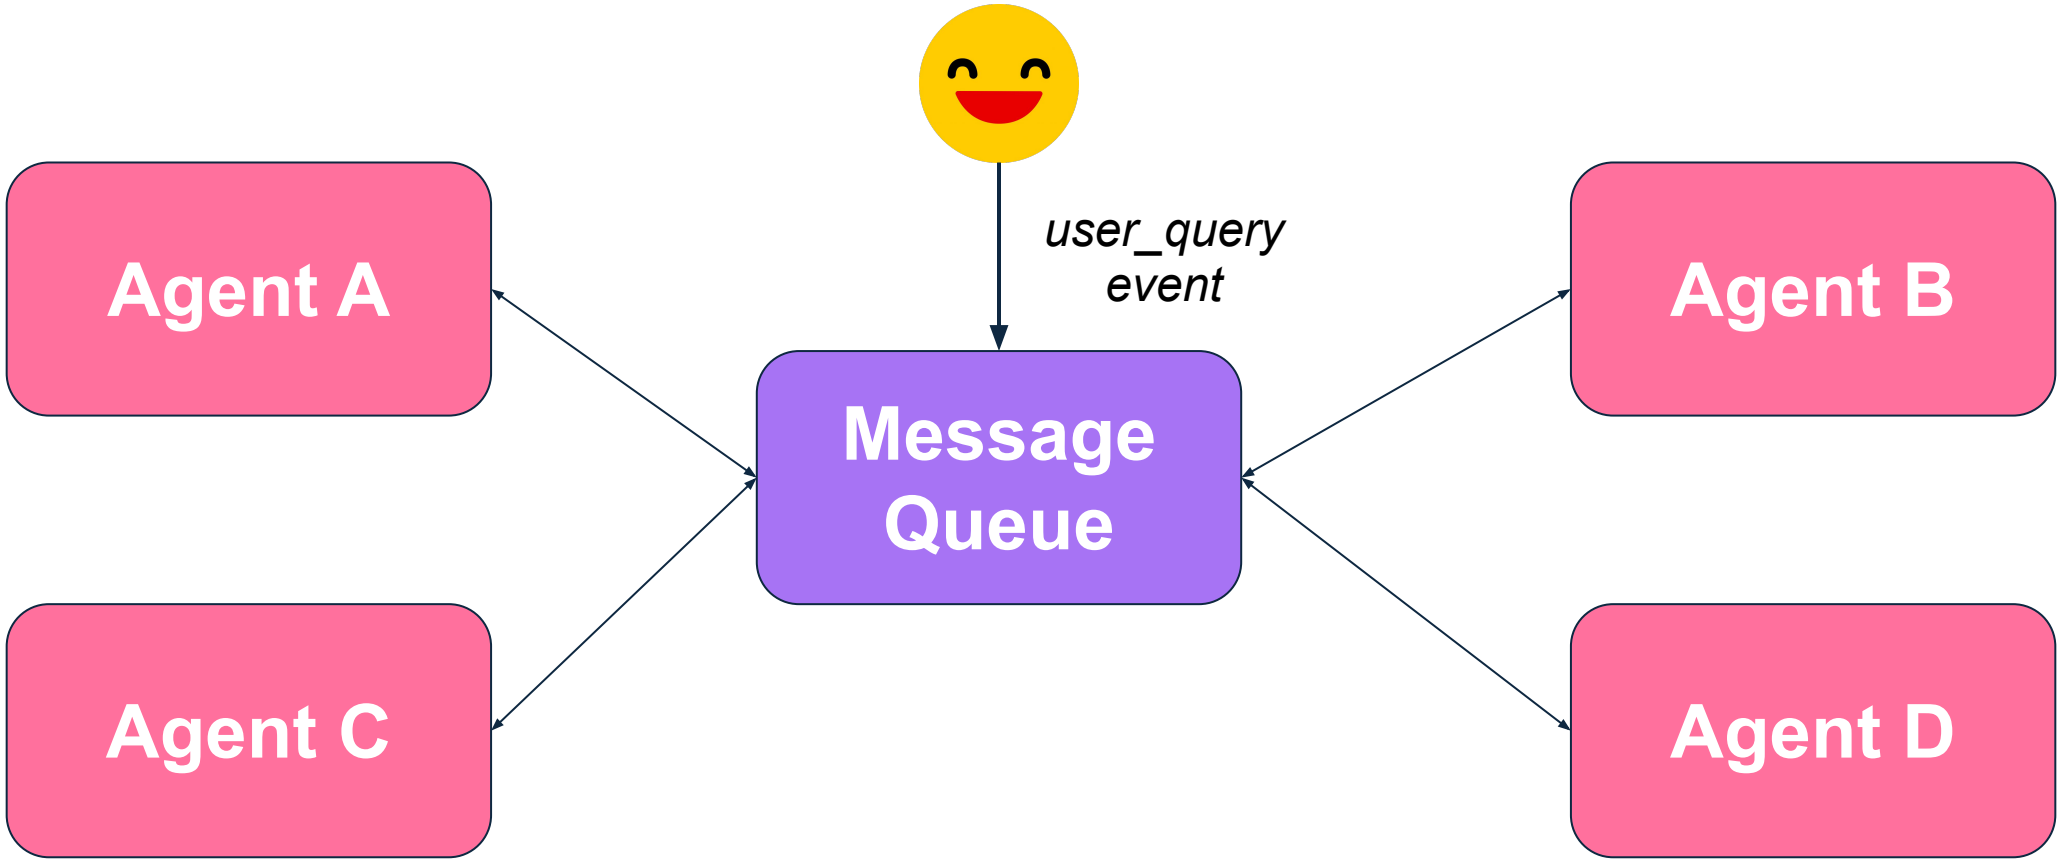

NUNNARI
LABS

Top Down/ Supervisory

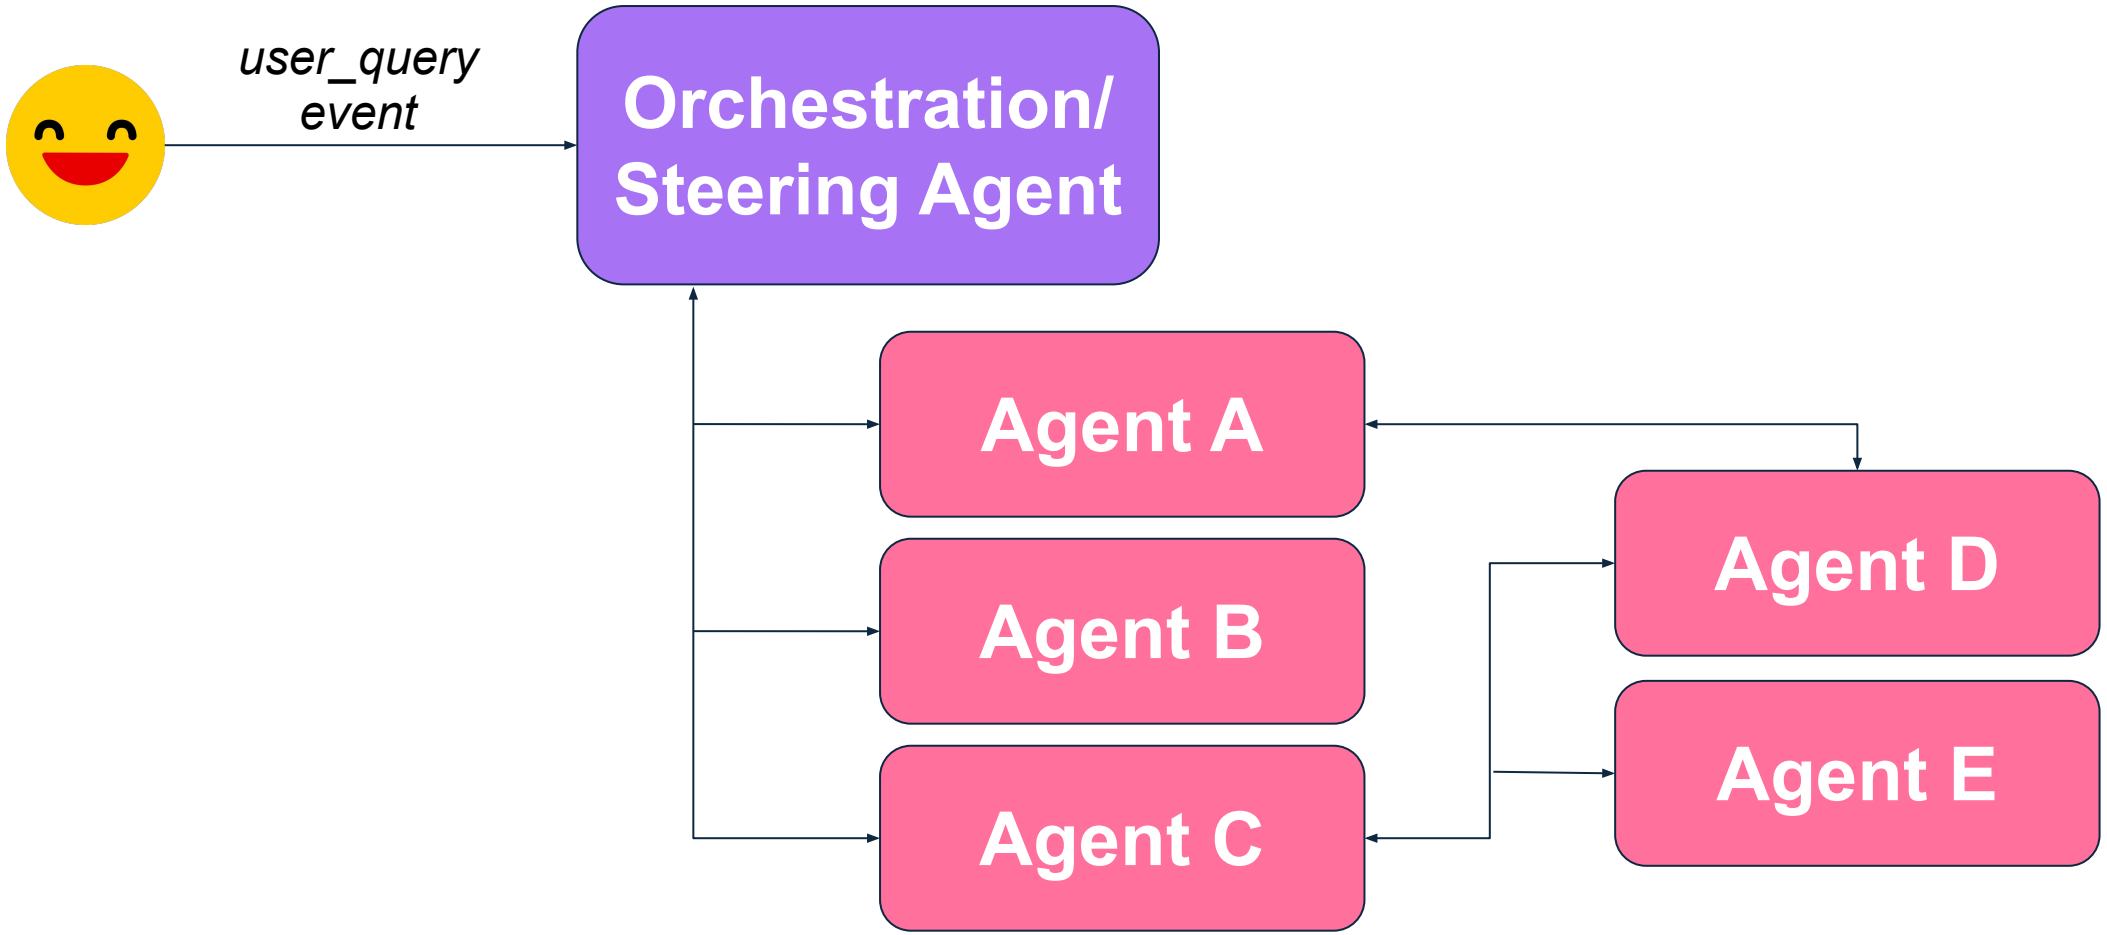

NUNNARI
LABS

AI Agent vs Multi-Agent System

Navaneeth Malingan

Aspect	Agent	Multi-Agent System
Definition	Individual AI application that achieves goals by observing and acting upon the world	System of multiple agents working together to solve complex tasks
Components	Foundational model, tools, reasoning loop	Multiple individual agents, communication system (e.g., messaging queue)
Architecture	Single unit	Collaborative or Top-down/Supervisory
Complexity	Simpler, focused on specific tasks	More complex, can handle diverse or multi-step tasks
Decision Making	Independent	Coordinated or hierarchical
Specialization	Usually focused on a particular domain or task	Can involve multiple specialized agents for different aspects of a task
Scalability	Limited to individual capabilities	More scalable for complex problems
Communication	Primarily with user/environment	Inter-agent communication as well as with user/environment
Examples	Customer support agent, Service Center agent	Collaborative problem-solving system, Orchestrated workflow system
Tools Usage	Uses tools directly	Multiple agents can use different tools as needed
Typical Use Cases	Specific, focused tasks (e.g., booking a flight)	Complex, multi-step processes (e.g., comprehensive customer service)

NUNNARI
LABS

Multi-Agentic Frameworks

Navaneeth Malingan

NUNNARI
LABS

Popular Frameworks for Building Agents

LangChain

LangGraph

LangGraph

NUNNARI LABS

LangGraph is a library for building **stateful**, **multi-actor** applications with **LLMs**, used to **create agent** and **multi-agent workflows**.

Key Advantages over other LLM frameworks

- **Cycles:** Unlike DAG-based solutions, LangGraph supports flows with cycles, crucial for developing complex agentic architectures.
- **Controllability:** Offers enhanced control over workflow processes, ensuring precise and adaptable operations.
- **Persistence:** Built-in persistence allows for advanced features such as human-in-the-loop interactions and memory retention, enabling more dynamic and responsive applications.

LangGraph - Key Features

NUNNARI LABS

- **Cycles and Branching**: Implement loops and conditionals in your apps.
- **Persistence**: Automatically save state after each step in the graph. Pause and resume the graph execution at any point to support error recovery, human-in-the-loop workflows, time travel and more.
- **Human-in-the-Loop**: Interrupt graph execution to approve or edit next action planned by the agent.
- **Streaming Support**: Stream outputs as they are produced by each node (including token streaming).
- **Integration with LangChain**: LangGraph integrates seamlessly with *LangChain* and *LangSmith* (but does not require them).

Why LangGraph?

NUNNARI LABS

Simplifies Development

a. State Management
b. Agent Coordination (Communication between Agents)
c. To Define Workflows and Logics

Flexibility

a. Flexibility to define custom agent logic and communication protocols.
b. Build tailored applications, from chatbots to complex multi-agent systems.
c. Provides the tools for highly customized solutions specific to your needs.

Why LangGraph?

NUNNARI LABS

Scalability

a. Large Scale Multi Agent Application
b. Handle high volume of Interaction and complex workflows
c. Enterprise level application (Cloud with drag and drop UI)

Fault Tolerance

a. Handle Error
b. It will not stop the flow when a agent fails

It’s Time to Code!

Navaneeth Malingan

References

- **All you need to know to Develop using Large Language Models**, Sergei Savvov - towardsdatascience.com
- **Fixing Hallucinations in LLMs**, Sergei Savvov - betterprogramming.pub
- **Intro to LLM Agents with LangChain: When RAG is Not Enough**, Alex Honchar - medium.com/towards-data-science

Thank You!

Navaneeth Malingan
navaneeth@nunnarilabs.com
www.nunnarilabs.com

NUNNARI LABS

AI Tamil Nadu

TensorFlow User Group Coimbatore

Connect with Me!